

Quantum Computing Lectures for Nano and Quantum Workshop

Morten Hjorth-Jensen^{1,2}

¹Department of Physics and Center for Computing in Science Education, University of Oslo, Norway

²Department of Physics and Astronomy and Facility for Rare Isotope Beams, Michigan State University, East Lansing, Michigan, USA

Cali, Colombia, December 4-8, 2023

Algorithms for solving quantum mechanical problems, quantum gates and circuits

1. Introduction to [Qiskit](#)
 2. Defining one-, two- and three-qubit gates
 3. Setting up quantum circuits and simple algorithms
- [Video of lecture](#)

Quantum gates, reminder from last week

Introduction to [Qiskit](#)

```
#!/usr/bin/env python
# coding: utf-8
import numpy as np
import qiskit as qk
from scipy.optimize import minimize

# # Initialize registers and circuit

n_qubits = 1 #Number of qubits
n_cbits = 1 #Number of classical bits (the number of qubits you want to measure at the end of the
qreg = qk.QuantumRegister(n_qubits) #Create a quantum register
creg = qk.ClassicalRegister(n_cbits) #Create a classical register
circuit = qk.QuantumCircuit(qreg,creg) #Create your quantum circuit

circuit.draw() #Draw circuit. It is empty
```

Thereafter we perform operations on qubit

```
circuit.x(qreg[0]) #Applies a Pauli X gate to the first qubit in the quantum register
circuit.draw()
```

and select a qubit to measure and encode the results to a classical bit

```
#Measure the first qubit in the quantum register
#and encode the results to the first qubit in the classical register
circuit.measure(qreg[0],creg[0])
circuit.draw()
```

Thereafter we execute the circuit

```
backend = qk.Aer.get_backend('qasm_simulator')
#This is the device you want to use. It is an ideal simulation of a quantum device
```

```
job = backend.run(circuit,shots=1000) #Run the circuit 1000 times
result = job.result()
counts = result.get_counts()
print(counts)
circuit.clear()
circuit.draw()
```

```
circuit.h(qreg[0]) #Apply a Hadamard gate to the first qubit of the quantum register
circuit.measure(qreg,creg)
print(circuit.draw())
```

```
job = backend.run(circuit,shots=1000)
result = job.result()
counts = result.get_counts()
print(counts)
circuit.clear()
```

Now we create a two-qubit circuit and set up a Bell state

```
n_qubits = 2
n_cbits = 2
qreg = qk.QuantumRegister(n_qubits)
creg = qk.ClassicalRegister(n_cbits)
circuit = qk.QuantumCircuit(qreg,creg)
circuit.draw()
```

```
circuit.h(qreg[0])
circuit.cx(qreg[0],qreg[1])
#This is a controlled operation. Apply a Pauli X gate to the second qubit (qreg[1]) if the first qubit
#is in the |1> state. Else do nothing

circuit.draw()

circuit.measure(qreg,creg)
circuit.draw()
```

```

job = backend.run(circuit,shots=1000)
result = job.result()
counts = result.get_counts()
print(counts)
circuit.clear()

```

We apply a rotation to a qubit

```

theta = np.pi/3
circuit.rx(theta, qreg[0]) #R_x(theta) rotation on the first qubit (qreg[0])
circuit.measure(qreg,creg)
print(circuit.draw())
job = backend.run(circuit,shots=1000)
result = job.result()
counts = result.get_counts()
circuit.clear()
print(counts)

```

Now we want to find the lowest eigenvalue of

$$H = c_1 Z_0 + c_2 Z_1 + c_3 X_0 Y_1$$

We will use

$$\langle \psi | H | \psi \rangle = c_1 \langle \psi | Z_0 | \psi \rangle + c_2 \langle \psi | Z_1 | \psi \rangle + c_3 \langle \psi | X_0 Y_1 | \psi \rangle$$

```

I = np.eye(2)
X = np.array([[0,1],[1,0]])
Y = np.array([[0,-1j],[1j,0]])
Z = np.array([[1,0],[0,-1]])
H = np.kron(Z,I) + np.kron(I,Z) + np.kron(X,Y)
eigvals,eigvecs = np.linalg.eigh(H)
print(eigvals[0])

```

```

c_1 = 1
c_2 = 1
c_3 = 1

```

```

h_1 = [c_1,[0],['z']]
h_2 = [c_2,[1],['z']]
h_3 = [c_3,[0,1],['x','y']]
H = [h_1,h_2,h_3]
H

```

```
H[0]
```

We create the ansatz

```

def ansatz(theta,n_qubits):
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)

```

```

        for i in range(n_qubits):
            circuit.ry(theta[i],qreg[i])
        for i in range(n_qubits-1):
            circuit.cx(qreg[i],qreg[i+1])
        return(circuit)
qreg = qk.QuantumRegister(n_qubits)
circuit = qk.QuantumCircuit(qreg)
circuit.h(qreg[:2])
print('Before ansatz')
print(circuit.draw())
theta = np.random.randn(2)
n_qubits = 2
circuit = circuit.compose(ansatz(theta,n_qubits))
print('After ansatz')
circuit.draw()

```

Now we change measurement basis

```

def basis_change(h_i,n_qubits):
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)

    for qubit,operator in zip(h_i[1],h_i[2]):
        if operator == 'x':
            circuit.h(qreg[qubit])
        if operator == 'y':
            circuit.sdg(qreg[qubit])
            circuit.h(qreg[qubit])
    return(circuit)
n_qubits = 2
qreg = qk.QuantumRegister(n_qubits)
circuit = qk.QuantumCircuit(qreg)
theta = np.random.randn(n_qubits)
circuit = circuit.compose(ansatz(theta,n_qubits))
print('Ansatz circuit')
circuit.draw()
circuit = circuit.compose(basis_change(H[2],n_qubits))
print('After basis transformation:')
print(circuit.draw())

```

Get energy for given rotational parameters θ

```

def get_energy(theta):
    n_qubits = 2
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)
    circuit = circuit.compose(ansatz(theta,n_qubits))
    circuit_list = []
    for idx,h_i in enumerate(H):
        basis_change_circuit = basis_change(h_i,n_qubits)
        new_circuit = circuit.compose(basis_change_circuit)
        creg = qk.ClassicalRegister(len(h_i[1]))
        new_circuit.add_register(creg)
        new_circuit.measure(qreg[h_i[1]],creg)
        circuit_list.append(new_circuit)
    shots = 10000
    job = backend.run(circuit_list,shots=shots)
    E = np.zeros(len(circuit_list))

```

```

for i in range(len(circuit_list)):
    result = job.result()
    counts = result.get_counts(i)
    for key,value in counts.items():
        e = 1
        for bit in key:
            if bit == '0':
                e *= 1
            if bit == '1':
                e *= -1
        E[i] += e*value
    E[i] *= H[i][0]
E /= shots
return(np.sum(E))

theta = np.random.randn(2)
get_energy(theta)

```

Minimize energy with Scipy

```

theta = np.random.randn(2)
res = minimize(get_energy, theta, method='Powell',tol=1e-12)
get_energy(res.x)

```

We define a more flexible ansatz

```

def ansatz(theta,n_qubits):
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)
    idx = 0
    for i in range(n_qubits):
        circuit.ry(theta[idx],qreg[i])
        idx += 1
    for i in range(n_qubits-1):
        circuit.cx(qreg[i],qreg[i+1])
    for i in range(n_qubits):
        circuit.rx(theta[idx],qreg[i])
        idx += 1
    for i in range(n_qubits-1):
        circuit.cx(qreg[i],qreg[i+1])
    return(circuit)
theta = np.random.randn(4)
res = minimize(get_energy, theta, method='Powell',tol=1e-16)
get_energy(res.x)

```

Minimize energy with gradient descent

$$\frac{\partial E(\theta_1, \dots, \theta_i, \dots, \theta_p)}{\partial \theta_i} = \frac{E(\theta_1, \dots, \theta_i + \pi/2, \dots, \theta_p) - E(\theta_1, \dots, \theta_i - \pi/2, \dots, \theta_p)}{2}$$

```

epochs = 200
theta = np.random.randn(4)
for epoch in range(epochs):
    print(epoch,get_energy(theta))
    grad = np.zeros_like(theta)
    for idx in range(theta.shape[0]):

```

```

theta_temp = theta.copy()
theta_temp[idx] += np.pi/2
E_plus = get_energy(theta_temp)
theta_temp[idx] -= np.pi
E_minus = get_energy(theta_temp)
grad[idx] = (E_plus - E_minus)/2
theta -= 0.1*grad

```

Solving quantum mechanical problems

1. Simple Hamiltonian, the Lipkin model
2. Introducing the Variational Quantum Eigensolver (VQE)

Simple Hamiltonian, the Lipkin model

We will study a schematic model (the Lipkin model, see Nuclear Physics **62** (1965) 188), for the interaction among 2 and more fermions that can occupy two different energy levels.

For four fermions, the case we consider first here, each levels has degeneration $d = 4$, leading to different total spin values. The two levels have quantum numbers $\sigma = \pm 1$, with the upper level having $2\sigma = +1$ and energy $\varepsilon_1 = \varepsilon/2$. The lower level has $2\sigma = -1$ and energy $\varepsilon_2 = -\varepsilon/2$. That is, the lowest single-particle level has negative spin projection (or spin down), while the upper level has spin up. In addition, the substates of each level are characterized by the quantum numbers $p = 1, 2, 3, 4$.

We define the single-particle states (for the four fermion case which we will work on here)

$$|u_{\sigma=-1,p}\rangle = a_{-p}^\dagger |0\rangle \quad |u_{\sigma=1,p}\rangle = a_{+p}^\dagger |0\rangle.$$

The single-particle states span an orthonormal basis. The Hamiltonian of the system is given by

$$\begin{aligned}
\hat{H} &= \hat{H}_0 + \hat{H}_1 + \hat{H}_2 \\
\hat{H}_0 &= \frac{1}{2}\varepsilon \sum_{\sigma,p} \sigma a_{\sigma,p}^\dagger a_{\sigma,p} \\
\hat{H}_1 &= \frac{1}{2}V \sum_{\sigma,p,p'} a_{\sigma,p}^\dagger a_{\sigma,p'}^\dagger a_{-\sigma,p'} a_{-\sigma,p} \\
\hat{H}_2 &= \frac{1}{2}W \sum_{\sigma,p,p'} a_{\sigma,p}^\dagger a_{-\sigma,p'}^\dagger a_{\sigma,p'} a_{-\sigma,p}
\end{aligned}$$

where V and W are constants. The operator H_1 can move pairs of fermions while H_2 is a spin-exchange term. The latter moves a pair of fermions from a state $(p\sigma, p' - \sigma)$ to a state $(p - \sigma, p'\sigma)$.

We are going to rewrite the above Hamiltonian in terms of so-called quasispin operators

$$\begin{aligned}\hat{J}_+ &= \sum_p a_{p+}^\dagger a_{p-} \\ \hat{J}_- &= \sum_p a_{p-}^\dagger a_{p+} \\ \hat{J}_z &= \frac{1}{2} \sum_{p\sigma} \sigma a_{p\sigma}^\dagger a_{p\sigma} \\ \hat{J}^2 &= J_+ J_- + J_z^2 - J_z\end{aligned}$$

We show here that these operators obey the commutation relations for angular momentum.

We can in turn express \hat{H} in terms of the above quasispin operators and the number operator

$$\hat{N} = \sum_{p\sigma} a_{p\sigma}^\dagger a_{p\sigma}.$$

We have the following quasispin operators

$$J_\pm = \sum_p a_{p\pm}^\dagger a_{p\mp}, \quad (1)$$

$$J_z = \frac{1}{2} \sum_{p,\sigma} \sigma a_{p\sigma}^\dagger a_{p\sigma}, \quad (2)$$

$$J^2 = J_+ J_- + J_z^2 - J_z, \quad (3)$$

and we want to compute the commutators

$$[J_z, J_\pm], \quad [J_+, J_-], \quad [J^2, J_\pm] \quad \text{og} \quad [J^2, J_z].$$

Let us start with the first one and inserting for J_z and J_\pm given by the equations (2) and (1), respectively, we obtain

$$\begin{aligned}[J_z, J_\pm] &= J_z J_\pm - J_\pm J_z \\ &= \left(\frac{1}{2} \sum_{p,\sigma} \sigma a_{p\sigma}^\dagger a_{p\sigma} \right) \left(\sum_{p'} a_{p'\pm}^\dagger a_{p'\mp} \right) - \left(\sum_{p'} a_{p'\pm}^\dagger a_{p'\mp} \right) \left(\frac{1}{2} \sum_{p,\sigma} \sigma a_{p\sigma}^\dagger a_{p\sigma} \right) \\ &= \frac{1}{2} \sum_{p,p',\sigma} \sigma \left(a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} - a_{p'\pm}^\dagger a_{p'\mp} a_{p\sigma}^\dagger a_{p\sigma} \right).\end{aligned}$$

Using the commutation relations for the creation and annihilation operators

$$\{a_l, a_k\} = 0, \quad (4)$$

$$\{a_l^\dagger, a_k^\dagger\} = 0, \quad (5)$$

$$\{a_l^\dagger, a_k\} = \delta_{lk}, \quad (6)$$

in order to move the operators in the right product to be in the same order as those in the lefthand product

$$\begin{aligned} [J_z, J_\pm] &= \frac{1}{2} \sum_{p,p',\sigma} \sigma \left(a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} - a_{p'\pm}^\dagger (\delta_{p'p} \delta_{\mp\sigma} - a_{p\sigma}^\dagger a_{p'\mp}) a_{p\sigma} \right) \\ &= \frac{1}{2} \sum_{p,p',\sigma} \sigma \left(a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} - a_{p'\pm}^\dagger \delta_{p'p} \delta_{\mp\sigma} a_{p\sigma} + a_{p'\pm}^\dagger a_{p\sigma}^\dagger a_{p'\mp} a_{p\sigma} \right), \end{aligned}$$

which results in

$$\begin{aligned} [J_z, J_\pm] &= \frac{1}{2} \sum_{p,p',\sigma} \sigma \left(a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} - a_{p'\pm}^\dagger \delta_{pp'} \delta_{\mp\sigma} a_{p\sigma} + a_{p\sigma}^\dagger a_{p'\pm}^\dagger a_{p\sigma} a_{p'\mp} \right) \\ &= \frac{1}{2} \sum_{p,p',\sigma} \sigma \left(a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} - a_{p'\pm}^\dagger \delta_{pp'} \delta_{\mp\sigma} a_{p\sigma} + a_{p\sigma}^\dagger (\delta_{pp'} \delta_{\pm\sigma} - a_{p\sigma} a_{p'\pm}^\dagger) a_{p'\mp} \right) \\ &= \frac{1}{2} \sum_{p,p',\sigma} \sigma \left(a_{p\sigma}^\dagger \delta_{pp'} \delta_{\pm\sigma} a_{p'\mp} - a_{p'\pm}^\dagger \delta_{pp'} \delta_{\mp\sigma} a_{p\sigma} \right). \end{aligned}$$

The last equality leads to

$$\begin{aligned} [J_z, J_\pm] &= \frac{1}{2} \sum_p \left((\pm 1) a_{p\pm}^\dagger a_{p\mp} - (\mp 1) a_{p\pm}^\dagger a_{p\mp} \right) = \pm \frac{1}{2} \sum_p \left(a_{p\pm}^\dagger a_{p\mp} + (\pm 1) a_{p\pm}^\dagger a_{p\mp} \right) \\ &= \pm \sum_p a_{p\pm}^\dagger a_{p\mp} = \pm J_\pm, \end{aligned}$$

where the last results follows from comparing with Eq. (1).

We can then continue with the next commutation relation, using Eq. (1),

$$\begin{aligned}
[J_+, J_-] &= J_+ J_- - J_- J_+ \\
&= \left(\sum_p a_{p'+}^\dagger a_{p-} \right) \left(\sum_{p'} a_{p'-}^\dagger a_{p'+} \right) - \left(\sum_{p'} a_{p'-}^\dagger a_{p'+} \right) \left(\sum_p a_{p'+}^\dagger a_{p-} \right) \\
&= \sum_{p,p'} \left(a_{p'+}^\dagger a_{p-} a_{p'-}^\dagger a_{p'+} - a_{p'-}^\dagger a_{p'+} a_{p'+}^\dagger a_{p-} \right) \\
&= \sum_{p,p'} \left(a_{p'+}^\dagger a_{p-} a_{p'-}^\dagger a_{p'+} - a_{p'-}^\dagger \left(\delta_{++} \delta_{pp'} - a_{p+}^\dagger a_{p'+} \right) a_{p-} \right) \\
&= \sum_{p,p'} \left(a_{p'+}^\dagger a_{p-} a_{p'-}^\dagger a_{p'+} - a_{p'-}^\dagger \delta_{pp'} a_{p-} + a_{p'+}^\dagger a_{p+}^\dagger a_{p'+} a_{p-} \right) \\
&= \sum_{p,p'} \left(a_{p'+}^\dagger a_{p-} a_{p'-}^\dagger a_{p'+} - a_{p'-}^\dagger \delta_{pp'} a_{p-} + a_{p+}^\dagger a_{p'+}^\dagger a_{p-} a_{p'+} \right) \\
&= \sum_{p,p'} \left(a_{p'+}^\dagger a_{p-} a_{p'-}^\dagger a_{p'+} - a_{p'-}^\dagger \delta_{pp'} a_{p-} + a_{p+}^\dagger \left(\delta_{--} \delta_{pp'} - a_{p-} a_{p'-}^\dagger \right) a_{p'+} \right) \\
&= \sum_{p,p'} \left(a_{p+}^\dagger \delta_{pp'} a_{p'+} - a_{p'-}^\dagger \delta_{pp'} a_{p-} \right),
\end{aligned}$$

which results in

$$[J_+, J_-] = \sum_p \left(a_{p+}^\dagger a_{p+} - a_{p-}^\dagger a_{p-} \right) = 2J_z,$$

It is straightforward to show that

$$[J^2, J_\pm] = [J_+ J_- + J_z^2 - J_z, J_\pm] = [J_+ J_-, J_\pm] + [J_z^2, J_\pm] - [J_z, J_\pm].$$

Using the relations

$$[AB, C] = A[B, C] + [A, C]B, \quad (7)$$

$$[A, BC] = [A, B]C + B[A, C], \quad (8)$$

we obtain

$$[J^2, J_\pm] = J_+[J_-, J_\pm] + [J_+, J_\pm]J_- + J_z[J_z, J_\pm] + [J_z, J_\pm]J_z - [J_z, J_\pm].$$

Finally, from the above it follows that

$$\begin{aligned}
[J^2, J_+] &= -2J_+ J_z + J_z[J_z, J_+] + [J_z, J_+]J_z - [J_z, J_+] \\
&= -2J_+ J_z + J_z J_+ + J_+ J_z - J_+ \\
&= -2J_+ J_z + J_+ + J_+ J_z + J_+ J_z - J_+ = 0,
\end{aligned}$$

and

$$\begin{aligned}
[J^2, J_-] &= 2J_z J_- + J_z[J_z, J_-] + [J_z, J_-]J_z - [J_z, J_-] \\
&= 2J_z J_- - J_z J_- - J_- J_z + J_- \\
&= J_z J_- - (J_z J_- + J_-) + J_- = 0.
\end{aligned}$$

Our last commutator is given by

$$\begin{aligned}
[J^2, J_z] &= [J_+ J_- + J_z^2 - J_z, J_z] \\
&= [J_+ J_-, J_z] + [J_z^2, J_z] - [J_z, J_z] \\
&= J_+[J_-, J_z] + [J_+, J_z]J_- \\
&= J_+ J_- - J_+ J_- = 0
\end{aligned}$$

Summing up we have

$$[J_z, J_\pm] = \pm J_\pm, \quad (9)$$

$$[J_+, J_-] = 2J_z, \quad (10)$$

$$[J^2, J_\pm] = 0, \quad (11)$$

$$[J^2, J_z] = 0, \quad (12)$$

which are the standard commutation relations for angular (or orbital) momentum L_\pm , L_z og L^2 .

Rewriting the Hamiltonian

We wrote the above Hamiltonian as

$$H = H_0 + H_1 + H_2,$$

with

$$H_0 = \frac{1}{2}\varepsilon \sum_{p\sigma} \sigma a_{p\sigma}^\dagger a_{p\sigma},$$

and

$$H_1 = \frac{1}{2}V \sum_{p,p',\sigma} a_{p\sigma}^\dagger a_{p'\sigma}^\dagger a_{p'-\sigma} a_{p-\sigma},$$

and

$$H_2 = \frac{1}{2}W \sum_{p,p',\sigma} a_{p\sigma}^\dagger a_{p'-\sigma}^\dagger a_{p'\sigma} a_{p-\sigma}.$$

We will now rewrite the Hamiltonian in terms of the above quasi-spin operators and the number operator

$$N = \sum_{p,\sigma} a_{p\sigma}^\dagger a_{p\sigma}. \quad (13)$$

Going through each term of the Hamiltonian and using the expressions for the quasi-spin operators we obtain

$$H_0 = \varepsilon J_z. \quad (14)$$

Moving over to H_1 and using the anti-commutation relations (4) through (6) we obtain

$$\begin{aligned} H_1 &= \frac{1}{2}V \sum_{p,p',\sigma} a_{p\sigma}^\dagger a_{p'\sigma}^\dagger a_{p'-\sigma} a_{p-\sigma} \\ &= \frac{1}{2}V \sum_{p,p',\sigma} -a_{p\sigma}^\dagger a_{p'\sigma}^\dagger a_{p-\sigma} a_{p'-\sigma} \\ &= \frac{1}{2}V \sum_{p,p',\sigma} -a_{p\sigma}^\dagger \left(\delta_{pp'} \delta_{\sigma-\sigma} - a_{p-\sigma} a_{p'\sigma}^\dagger \right) a_{p'-\sigma} \\ &= \frac{1}{2}V \sum_{p,p',\sigma} a_{p\sigma}^\dagger a_{p-\sigma} a_{p'\sigma}^\dagger a_{p'-\sigma} \end{aligned}$$

Rewriting the sum over σ we arrive at

$$\begin{aligned} H_1 &= \frac{1}{2}V \sum_{p,p'} a_{p+}^\dagger a_{p-} a_{p'+}^\dagger a_{p'-} + a_{p-}^\dagger a_{p+} a_{p'-}^\dagger a_{p'+} \\ &= \frac{1}{2}V \left[\sum_p \left(a_{p+}^\dagger a_{p-} \right) \sum_{p'} \left(a_{p'+}^\dagger a_{p'-} \right) + \sum_p \left(a_{p-}^\dagger a_{p+} \right) \sum_{p'} \left(a_{p'-}^\dagger a_{p'+} \right) \right] \\ &= \frac{1}{2}V [J_+ J_+ + J_- J_-] = \frac{1}{2}V [J_+^2 + J_-^2], \end{aligned}$$

which leads to

$$H_1 = \frac{1}{2}V (J_+^2 + J_-^2). \quad (15)$$

Finally, we rewrite the last term

$$\begin{aligned}
H_2 &= \frac{1}{2}W \sum_{p,p',\sigma} a_{p\sigma}^\dagger a_{p'-\sigma}^\dagger a_{p'\sigma} a_{p-\sigma} \\
&= \frac{1}{2}W \sum_{p,p',\sigma} -a_{p\sigma}^\dagger a_{p'-\sigma}^\dagger a_{p-\sigma} a_{p'\sigma} \\
&= \frac{1}{2}W \sum_{p,p',\sigma} -a_{p\sigma}^\dagger \left(\delta_{pp'} \delta_{-\sigma-\sigma} - a_{p-\sigma} a_{p'-\sigma}^\dagger \right) a_{p'\sigma} \\
&= \frac{1}{2}W \sum_{p,p',\sigma} -a_{p\sigma}^\dagger \delta_{pp'} a_{p'\sigma} + a_{p\sigma}^\dagger a_{p-\sigma} a_{p'-\sigma}^\dagger a_{p'\sigma} \\
&= \frac{1}{2}W \left(-\sum_{p,\sigma} a_{p\sigma}^\dagger a_{p\sigma} + \sum_{p,p',\sigma} a_{p\sigma}^\dagger a_{p-\sigma} a_{p'-\sigma}^\dagger a_{p'\sigma} \right)
\end{aligned}$$

Using the expression for the number operator we obtain

$$\begin{aligned}
\sum_{p,p',\sigma} a_{p\sigma}^\dagger a_{p-\sigma} a_{p'-\sigma}^\dagger a_{p'\sigma} &= \sum_{p,p'} a_{p+}^\dagger a_{p-} a_{p'-}^\dagger a_{p'+} + a_{p-}^\dagger a_{p+} a_{p'+}^\dagger a_{p'-} \\
&= \sum_p \left(a_{p+}^\dagger a_{p-} \right) \sum_{p'} \left(a_{p'-}^\dagger a_{p'+} \right) + \sum_p \left(a_{p-}^\dagger a_{p+} \right) \sum_{p'} \left(a_{p'+}^\dagger a_{p'-} \right) \\
&= J_+ J_- + J_- J_+,
\end{aligned}$$

resulting in

$$H_2 = \frac{1}{2}W (-N + J_+ J_- + J_- J_+). \quad (16)$$

We have thus expressed the Hamiltonian in term of the quasi-spin operators.

Commutation relations for the Hamiltonian

The above expressions can in turn be used to show that the Hamiltonian commutes with the various quasi-spin operators. This leads to quantum numbers which are conserved. Let us first show that $[H, J^2] = 0$, which means that J is a so-called *good* quantum number and that the total spin is a conserved quantum number.

We have

$$\begin{aligned}
[H, J^2] &= [H_0 + H_1 + H_2, J^2] \\
&= [H_0, J^2] + [H_1, J^2] + [H_2, J^2] \\
&= \varepsilon[J_z, J^2] + \frac{1}{2}V[J_+^2 + J_-^2, J^2] + \frac{1}{2}W[-N + J_+ J_- + J_- J_+, J^2].
\end{aligned}$$

We have previously shown that

$$[H, J^2] = \frac{1}{2}V ([J_+^2, J^2] + [J_-^2, J^2]) + \frac{1}{2}W (-[N, J^2] + [J_+J_-, J^2] + [J_-J_+, J^2])$$

Using that $[J_\pm, J^2] = 0$, it follows that $[J_\pm^2, J^2] = 0$. We can then see that $[J_+J_-, J^2] = 0$ and $[J_-J_+, J^2] = 0$ which leads to

$$\begin{aligned} [H, J^2] &= -\frac{1}{2}W[N, J^2] \\ &= \frac{1}{2}W (-[N, J_+J_-] - [N, J_z^2] + [N, J_z]) \\ &= \frac{1}{2}W (-[N, J_+]J_- - J_+[N, J_-] - [N, J_z]J_z - J_z[N, J_z] + [N, J_z]). \end{aligned}$$

Combining with the number operator we have

$$\begin{aligned} [N, J_\pm] &= NJ_\pm - J_\pm N \\ &= \left(\sum_{p, \sigma} a_{p\sigma}^\dagger a_{p\sigma} \right) \left(\sum_{p'} a_{p'\pm}^\dagger a_{p'\mp} \right) - \left(\sum_{p'} a_{p'\pm}^\dagger a_{p'\mp} \right) \left(\sum_{p, \sigma} a_{p\sigma}^\dagger a_{p\sigma} \right) \\ &= \sum_{p, p', \sigma} a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} - a_{p'\pm}^\dagger a_{p'\mp} a_{p\sigma}^\dagger a_{p\sigma} \\ &= \sum_{p, p', \sigma} a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} - a_{p'\pm}^\dagger (\delta_{\mp\sigma} \delta_{pp'} - a_{p\sigma}^\dagger a_{p'\mp}) a_{p\sigma} \\ &= \sum_{p, p', \sigma} a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} - a_{p'\pm}^\dagger \delta_{\mp\sigma} \delta_{pp'} a_{p\sigma} + a_{p'\pm}^\dagger a_{p\sigma}^\dagger a_{p'\mp} a_{p\sigma} \\ &= \sum_{p, p', \sigma} a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} + a_{p\sigma}^\dagger a_{p'\pm}^\dagger a_{p\sigma} a_{p'\mp} - \sum_p a_{p\pm}^\dagger a_{p\mp} \\ &= \sum_{p, p', \sigma} a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} + a_{p\sigma}^\dagger (\delta_{pp'} \delta_{\pm\sigma} - a_{p\sigma} a_{p'\pm}^\dagger) a_{p'\mp} - \sum_p a_{p\pm}^\dagger a_{p\mp} \\ &= \sum_p a_{p\pm}^\dagger a_{p\mp} - \sum_p a_{p\pm}^\dagger a_{p\mp} = 0. \end{aligned}$$

We obtain then

$$\begin{aligned} [N, J_z] &= NJ_z - J_z N \\ &= \left(\sum_{p, \sigma} a_{p\sigma}^\dagger a_{p\sigma} \right) \left(\frac{1}{2} \sum_{p', \sigma} \sigma a_{p'\sigma}^\dagger a_{p'\sigma} \right) - \left(\frac{1}{2} \sum_{p', \sigma} \sigma a_{p'\sigma}^\dagger a_{p'\sigma} \right) \left(\sum_{p, \sigma} a_{p\sigma}^\dagger a_{p\sigma} \right) \\ &= \sum_{p, p', \sigma} \sigma a_{p\sigma}^\dagger a_{p\sigma} a_{p'\sigma}^\dagger a_{p'\sigma} - \sigma a_{p'\sigma}^\dagger a_{p'\sigma} a_{p\sigma}^\dagger a_{p\sigma} = 0, \end{aligned}$$

which leads to

$$[H, J^2] = 0, \tag{17}$$

and J is a good quantum number.

Constructing the Hamiltonian matrix for $J = 2$. We start with the state (unique) where all spins point down

$$|2, -2\rangle = a_{1-}^\dagger a_{2-}^\dagger a_{3-}^\dagger a_{4-}^\dagger |0\rangle \quad (18)$$

which is a state with $J_z = -2$ and $J = 2$. (we label the states as $|J, J_z\rangle$). For $J = 2$ we have the spin projections $J_z = -2, -1, 0, 1, 2$. We can use the lowering and raising operators for spin in order to define the other states

$$J_+ |J, J_z\rangle = \sqrt{J(J+1) - J_z(J_z+1)} |J, J_z+1\rangle, \quad (19)$$

$$J_- |J, J_z\rangle = \sqrt{J(J+1) - J_z(J_z-1)} |J, J_z-1\rangle. \quad (20)$$

We can then construct all other states with $J = 2$ using the raising operator J_+ on $|2, -2\rangle$

$$J_+ |2, -2\rangle = \sqrt{2(2+1) - (-2)(-2+1)} |2, -2+1\rangle = \sqrt{6-2} |2, -1\rangle = 2 |2, -1\rangle,$$

which gives

$$\begin{aligned} |2, -1\rangle &= \frac{1}{2} J_+ |2, -2\rangle \\ &= \frac{1}{2} \sum_p a_{p+}^\dagger a_{p-} a_{1-}^\dagger a_{2-}^\dagger a_{3-}^\dagger a_{4-}^\dagger |0\rangle \\ &= \frac{1}{2} \left(a_{1+}^\dagger a_{2-}^\dagger a_{3-}^\dagger a_{4-}^\dagger + a_{1-}^\dagger a_{2+}^\dagger a_{3-}^\dagger a_{4-}^\dagger + a_{1-}^\dagger a_{2-}^\dagger a_{3+}^\dagger a_{4-}^\dagger + a_{1-}^\dagger a_{2-}^\dagger a_{3-}^\dagger a_{4+}^\dagger \right) |0\rangle. \end{aligned} \quad (21)$$

We can construct all the other states in the same way. That is

$$J_+ |2, -1\rangle = \sqrt{2(2+1) - (-1)(-1+1)} |2, -1+1\rangle = \sqrt{6} |2, 0\rangle,$$

which results in

$$\begin{aligned} |2, 0\rangle &= \frac{1}{\sqrt{6}} \left(a_{1+}^\dagger a_{2+}^\dagger a_{3-}^\dagger a_{4-}^\dagger + a_{1+}^\dagger a_{2-}^\dagger a_{3+}^\dagger a_{4-}^\dagger + a_{1+}^\dagger a_{2-}^\dagger a_{3-}^\dagger a_{4+}^\dagger + a_{1-}^\dagger a_{2+}^\dagger a_{3+}^\dagger a_{4-}^\dagger + \right. \\ &\quad \left. a_{1-}^\dagger a_{2+}^\dagger a_{3-}^\dagger a_{4+}^\dagger + a_{1-}^\dagger a_{2-}^\dagger a_{3+}^\dagger a_{4+}^\dagger \right) |0\rangle \end{aligned} \quad (22)$$

The two remaining states are

$$|2, 1\rangle = \frac{1}{2} \left(a_{1+}^\dagger a_{2+}^\dagger a_{3+}^\dagger a_{4-}^\dagger + a_{1+}^\dagger a_{2+}^\dagger a_{3-}^\dagger a_{4+}^\dagger + a_{1+}^\dagger a_{2-}^\dagger a_{3+}^\dagger a_{4+}^\dagger + a_{1-}^\dagger a_{2+}^\dagger a_{3+}^\dagger a_{4+}^\dagger \right). \quad (23)$$

and

$$|2, 2\rangle = a_{1+}^\dagger a_{2+}^\dagger a_{3+}^\dagger a_{4+}^\dagger |0\rangle. \quad (24)$$

These five states can in turn be used as computational basis states in order to define the Hamiltonian matrix to be diagonalized. The matrix elements are

given by $\langle J, J_z | H | J', J'_z \rangle$. The Hamiltonian is hermitian and we obtain after all this labor of ours

$$H_{J=2} = \begin{bmatrix} -2\varepsilon & 0 & \sqrt{6}V & 0 & 0 \\ 0 & -\varepsilon + 3W & 0 & 3V & 0 \\ \sqrt{6}V & 0 & 4W & 0 & \sqrt{6}V \\ 0 & 3V & 0 & \varepsilon + 3W & 0 \\ 0 & 0 & \sqrt{6}V & 0 & 2\varepsilon \end{bmatrix} \quad (25)$$

We can now select a set of parameters and diagonalize the above matrix. We select $\varepsilon = 2$, $V = -1/3$, $W = -1/4$ and our matrix becoes

$$H_{J=2}^{(1)} = \begin{bmatrix} -4 & 0 & -\sqrt{6}/3 & 0 & 0 \\ 0 & -2 - 3/4 & 0 & -1 & 0 \\ -\sqrt{6}/3 & 0 & -1 & 0 & -\sqrt{6}/3 \\ 0 & -1 & 0 & 2 + -3/4 & 0 \\ 0 & 0 & -\sqrt{6}/3 & 0 & 4 \end{bmatrix},$$

which gives the eigenvalue

$$D = \begin{bmatrix} -4.21288 & 0 & 0 & 0 & 0 \\ 0 & -2.98607 & 0 & 0 & 0 \\ 0 & 0 & -0.91914 & 0 & 0 \\ 0 & 0 & 0 & 1.48607 & 0 \\ 0 & 0 & 0 & 0 & 4.13201 \end{bmatrix}.$$

The lowest state has an admixture of basis states given by

$$|\psi_0\rangle = 0.96735|2, -2\rangle + 0.25221|2, 0\rangle + 0.02507|2, 2\rangle,$$

with energy $E_0 = -4.21288$.

We can now change the parameters to $\varepsilon = 2$, $V = -4/3$, $W = -1$. Our matrix reads then

$$H_{J=2}^{(2)} = \begin{bmatrix} -4 & 0 & -4\sqrt{6}/3 & 0 & 0 \\ 0 & -5 & 0 & -4 & 0 \\ -4\sqrt{6}/3 & 0 & -4 & 0 & -4\sqrt{6}/3 \\ 0 & -4 & 0 & -1 & 0 \\ 0 & 0 & -4\sqrt{6}/3 & 0 & 4 \end{bmatrix},$$

with the following eigenvalues

$$D = \begin{bmatrix} -7.75122 & 0 & 0 & 0 & 0 \\ 0 & -7.47214 & 0 & 0 & 0 \\ 0 & 0 & -1.55581 & 0 & 0 \\ 0 & 0 & 0 & 1.47214 & 0 \\ 0 & 0 & 0 & 0 & 5.30704 \end{bmatrix}.$$

The new ground state (lowest state) has the following admixture of computational basis states

$$|\psi_0\rangle = 0.64268|2, -2\rangle + 0.73816|2, 0\rangle + 0.20515|2, 2\rangle,$$

with energy $E_0 = -7.75122$.

For the first set of parameters, the likelihood for observing the system in the computational basis state $|2, -2\rangle$ is rather large. This is expected since the interaction matrix elements are smaller than the single-particle energies. For the second case, with larger matrix elements, we see a much stronger mixing of the other states, again as expected due to the ratio of the interaction matrix elements and the single-particle energies.

To do: add simple code

Mean field solution, Hartree-Fock theory and comparison with exact diagonalization

We define now a new basis set (single-particle basis set) given by $|\phi_{\alpha,p}\rangle$ given by a unitary transformation of the original single-particle basis states $|u_{\sigma,p}\rangle$

$$|\phi_{\alpha,p}\rangle = \sum_{\sigma=\pm 1} C_{\alpha\sigma} |u_{\sigma,p}\rangle. \quad (26)$$

The new states can thus be written as a linear combination of the previous basis states (with fixed p). This means that for the new basis states, the quantum numbers p stay the same since they are just linear combinations of the previous basis states with the same value of p . Since the previous basis is defined to be orthogonal, it is straightforward to demonstrate that the new basis has to preserve the orthogonality since we performing a unitary transformation. We will not show that here.

A new Slater determinant can be constructed and is given by

$$|\Psi\rangle = \prod_{p=1}^4 b_{\alpha,p}^\dagger |0\rangle. \quad (27)$$

We can express the new creation (and annihilation operators) $b_{\alpha,p}^\dagger$ using $a_{p,\sigma}^\dagger$

$$|\phi_{\alpha,p}\rangle = b_{\alpha,p}^\dagger |0\rangle = \sum_{\sigma=\pm 1} C_{\alpha\sigma} |u_{\sigma,p}\rangle = \sum_{\sigma=\pm 1} C_{\alpha\sigma} a_{p\sigma}^\dagger |0\rangle = (C_{\alpha+} a_{p+}^\dagger + C_{\alpha-} a_{p-}^\dagger) |0\rangle.$$

With these expressions we can calculate the expectation value of the energy $E = \langle \Psi | H | \Psi \rangle$, resulting in

$$E = \langle \Psi | H | \Psi \rangle = \langle \Psi | (H_0 + H_1 + H_2) | \Psi \rangle,$$

which reads

$$E = 2\varepsilon |C_{\alpha+}|^2 - 2\varepsilon |C_{\alpha-}|^2 + \frac{1}{2} V (12(C_{\alpha+}^*)^2 (C_{\alpha-})^2 + 12(C_{\alpha+})^2 (C_{\alpha-}^*)^2) + 12W |C_{\alpha+}|^2 |C_{\alpha-}|^2.$$

If we assume that the coefficients $C_{\alpha\pm}$ are real, meaning that $C_{\alpha\pm}^* = C_{\alpha\pm}$, we obtain

$$E = 2\varepsilon C_{\alpha+}^2 - 2\varepsilon C_{\alpha-}^2 + 12(V + W)C_{\alpha+}^2 C_{\alpha-}^2.$$

We can now minimize the above equation with respect to the coefficients C and assuming they are real, we obtain (using that $C_{\alpha+}$ and $C_{\alpha-}$ are real and that the states $|\phi_{\alpha,p}\rangle$ are normalized, that $C_{\alpha+}^2 + C_{\alpha-}^2 = 1$). Inserted in the expression for the energy we obtain then

$$E = 2\varepsilon(1 - C_{\alpha-}^2) - 2\varepsilon C_{\alpha-}^2 + 12(V + W)(1 - C_{\alpha-}^2)C_{\alpha-}^2.$$

We simplify the above expression by defining $x = C_{\alpha-}^2$ og minimize the energy with respect to x

$$\begin{aligned} \frac{\partial E}{\partial x} &= \frac{\partial}{\partial x} [2\varepsilon(1 - x) - 2\varepsilon x + 12(V + W)(1 - x)x] \\ &= -2\varepsilon - 2\varepsilon + 12(V + W)(-x + 1 - x) \\ &= -4\varepsilon + 12(V + W)(1 - 2x). \end{aligned}$$

We obtain the minima by setting

$$\frac{\partial E}{\partial x} = 0 \quad \Rightarrow \quad -4\varepsilon + 12(V + W)(1 - 2x) = 0,$$

which results in

$$-\varepsilon + 3(V + W) = 6(V + W)x \quad \Rightarrow \quad x = \frac{1}{2} - \frac{\varepsilon}{6(V + W)}.$$

This gives the following values for $C_{\alpha+}$ and $C_{\alpha-}$:

$$C_{\alpha-}^2 = \frac{1}{2} - \frac{\varepsilon}{6(V + W)}, \quad C_{\alpha+}^2 = 1 - C_{\alpha-}^2 = \frac{1}{2} + \frac{\varepsilon}{6(V + W)}. \quad (28)$$

If we have found the energy minimum, this leads to the following stability criteria

$$(V + W) \leq \frac{\varepsilon}{3} \leq -(V + W), \quad (29)$$

where we have assumed that $V + W < 0$. If we have $V + W > 0$, we simply get reversed inequality signs.

Variational Quantum Eigensolver

One initial algorithm to estimate the eigenenergies of a quantum Hamiltonian was [quantum phase estimation](#). In it, one encodes the eigenenergies, one binary bit at a time (up to n bits), into the complex phases of the quantum states of the Hilbert space for n qubits. It does this by applying powers of controlled unitary evolution operators to a quantum state that can be expanded in terms of the Hamiltonian's eigenvectors of interest. The eigenenergies are encoded into the complex phases in such a way that taking the inverse quantum Fourier

transformation (see material on Quantum Fourier Transforms) of the states into which the eigen-energies are encoded results in a measurement probability distribution that has peaks around the bit strings that represent a binary fraction which corresponds to the eigen-energies of the quantum state acted upon by the controlled unitary operators. While quantum phase estimation (QPE) is provably efficient, non-hybrid, and non-variational, the number of qubits and length of circuits required is too great for our NISQ era quantum computers. Thus, QPE is only efficiently applicable to large, fault-tolerant quantum computers that likely won't exist in the near, but the far future.

Therefore, a different algorithm for finding the eigen-energies of a quantum Hamiltonian was put forth in 2014 called the variational quantum eigensolver, commonly referred to as VQE. The algorithm is hybrid, meaning that it requires the use of both a quantum computer and a classical computer. It is also variational, meaning that it relies, ultimately, on solving an optimization problem by varying parameters and thus is not as deterministic as QPE. The variational quantum eigensolver is based on the variational principle: The expectation value of a Hamiltonian H in a state $|\psi(\theta)\rangle$ parameterized by a set of angles θ , is always greater than or equal to the minimum eigen-energy E_0 . To see this, let $|n\rangle$ be the eigenstates of H , that is

$$H|n\rangle = E_n|n\rangle. \quad (30)$$

We can then expand our state $|\psi(\theta)\rangle$ in terms of said eigenstates

$$|\psi(\theta)\rangle = \sum_n c_n |n\rangle,$$

and plug this into the expectation value to yield

$$\langle\psi(\theta)|H|\psi(\theta)\rangle = \sum_{nm} c_m^* c_n \langle m|H|n\rangle = \sum_{nm} c_m^* c_n E_n \langle m|n\rangle = \sum_{nm} \delta_{nm} c_m^* c_n E_n = \sum_n |c_n|^2 E_n \geq E_0 \sum_n |c_n|^2 = E_0,$$

which implies that we can minimize over the set of angles θ and arrive at the ground state energy E_0

$$\min_{\theta} \langle\psi(\theta)|H|\psi(\theta)\rangle = E_0.$$

Using this fact, the VQE algorithm can be broken down into the following steps

1. Prepare the variational state $|\psi(\theta)\rangle$ on a quantum computer.
2. Measure this circuit in various bases and send these measurements to a classical computer
3. The classical computer post-processes the measurement data to compute the expectation value $\langle\psi(\theta)|H|\psi(\theta)\rangle$

4. The classical computer varies the parameters θ according to a classical minimization algorithm and sends them back to the quantum computer which runs step 1 again.

This loop continues until the classical optimization algorithm terminates which results in a set of angles θ_{\min} that characterize the ground state $|\phi(\theta_{\min})\rangle$ and an estimate for the ground state energy $\langle\psi(\theta_{\min})|H|\psi(\theta_{\min})\rangle$.

Expectation values

To execute the second step of VQE, we need to understand how expectation values of operators can be estimated via quantum computers by post-processing measurements of quantum circuits in different basis. To rotate bases, one uses the basis rotator R_σ which is defined for each Pauli gate σ to be

$$R_\sigma = H, \text{ if } \sigma = X, \quad (31)$$

and

$$HS^\dagger, \text{ if } \sigma = Y, \quad (32)$$

and

$$I, \text{ if } \sigma = Z. \quad (33)$$

We can show that these rotations allow us to measure the eigenvalues of the Pauli operators. The eigenvectors of the Pauli X gate are

$$|\pm\rangle = \frac{|0\rangle \pm |1\rangle}{\sqrt{2}},$$

with eigenvalues ± 1 . Acting on the eigenstates with the rotation in eq. (31) gives

$$H|+\rangle = +1|0\rangle,$$

and

$$H|-\rangle = -1|1\rangle.$$

Any single-qubit state can be written as a linear combination of these eigenvectors,

$$|\psi\rangle = \alpha|+\rangle + \beta|-\rangle.$$

We then have the following expectation value for the Pauli X operator

$$X = \psi X |\psi\rangle = |\alpha|^2 - |\beta|^2.$$

However, we can only measure the qubits in the computational basis. Applying the rotation in eq. (32) to our state gives

$$H|\psi\rangle = \alpha|0\rangle - \beta|1\rangle.$$

This tells us that we are able to estimate $|\alpha|^2$ and $|\beta|^2$ (and hence the expectation value of the Pauli X operator) by using the rotation in eq. () and measure the resulting state in the computational basis. We can show this for the Pauli Z and Pauli Y similarly.

Note the following identity of the basis rotator

$$R_\sigma^\dagger Z R_\sigma = \sigma,$$

which follows from the fact that $HZH = X$ and $SXS^\dagger = Y$. With this, we see that the expectation value of an arbitrary Pauli-gate σ in the state $|\psi\rangle$ can be expressed as a linear combination of probabilities

$$\begin{aligned} E_\psi(\sigma) &= \langle \psi | \sigma | \psi \rangle \\ &= \langle \psi | R_\sigma^\dagger Z R_\sigma | \psi \rangle = \langle \phi | Z | \phi \rangle \\ &= \langle \phi | \left(\sum_{x \in \{0,1\}} (-1)^x |x\rangle \langle x| \right) | \phi \rangle \\ &= \sum_{x \in \{0,1\}} (-1)^x |\langle x | \phi \rangle|^2 \\ &= \sum_{x \in \{0,1\}} (-1)^x P(|\phi\rangle \rightarrow |x\rangle), \end{aligned} \tag{34}$$

where $|\phi\rangle = |R_\sigma \psi\rangle$ and $P(|\phi\rangle \rightarrow |x\rangle)$ is the probability that the state $|\phi\rangle$ collapses to the state $|x\rangle$ when measured. This can be extended to any arbitrary Pauli string: consider the string of Pauli operators $P = \bigotimes_{p \in Q} \sigma_p$ which acts non-trivially on the set of qubits Q which is a subset of the total set of n qubits in the system. Then

$$\begin{aligned}
E_\psi(P) &= \langle \psi | \left(\bigotimes_{p \in Q} \sigma_p \right) | \psi \rangle \\
&= \langle \psi | \left(\bigotimes_{p \in Q} \sigma_p \right) \left(\bigotimes_{q \notin Q} I_q \right) | \psi \rangle \\
&= \langle \psi | \left(\bigotimes_{p \in Q} R_{\sigma_p}^\dagger Z_p R_{\sigma_p} \right) \left(\bigotimes_{q \notin Q} I_q \right) | \psi \rangle \\
&= \langle \psi | \left(\bigotimes_{p \in Q} R_{\sigma_p}^\dagger \right) \left(\bigotimes_{p \in Q} Z_p \right) \left(\bigotimes_{q \notin Q} I_q \right) \left(\bigotimes_{p \in Q} R_{\sigma_p} \right) | \psi \rangle \\
&= \langle \phi | \left(\bigotimes_{p \in Q} Z_p \right) \left(\bigotimes_{q \notin Q} I_q \right) | \phi \rangle \\
&= \langle \phi | \left(\bigotimes_{p \in Q} \sum_{x_p \in \{0,1\}} (-1)^{x_p} |x_p\rangle \langle x_p| \right) \left(\bigotimes_{q \notin Q} \sum_{y_q \in \{0,1\}} |y_q\rangle \langle y_q| \right) | \phi \rangle \\
&= \langle \phi | \left(\sum_{x \in \{0,1\}^n} (-1)^{\sum_{p \in Q} x_p} |x\rangle \langle x| \right) | \phi \rangle \\
&= \sum_{x \in \{0,1\}^n} (-1)^{\sum_{p \in Q} x_p} |\langle x | \phi \rangle|^2 \\
&= \sum_{x \in \{0,1\}^n} (-1)^{\sum_{p \in Q} x_p} P(|\phi\rangle \rightarrow |x\rangle), \tag{35}
\end{aligned}$$

where $|\phi\rangle = |\bigotimes_{p \in Q} R_{\sigma_p} \psi\rangle$. Finally, because the expectation value is linear

$$E_\psi \left(\sum_m \lambda_m P_m \right) = \sum_m \lambda_m E_\psi(P_m), \tag{36}$$

one can estimate any observable that can be written as a linear combination of Pauli-string terms.

Measurement

To estimate the probability $P(|\phi\rangle \rightarrow |x\rangle)$ from the previous section, one prepares the state $|\phi\rangle$ on a quantum computer and measures it, and then repeats this process (prepare and measure) several times. The probability $P(|\phi\rangle \rightarrow |x\rangle)$ is estimated to be the number of times that one measures the bit-string x divided by the total number of measurements that one makes; that is

$$P(|\phi\rangle \rightarrow |x\rangle) \approx \sum_{m=1}^M \frac{x_m}{M}, \quad (37)$$

where $x_m = 1$ if the result of measurement is x and 0 if the result of measurement is not x .

By the law of large numbers the approximation approaches equality as M goes to infinity

$$P(|\phi\rangle \rightarrow |x\rangle) = \lim_{M \rightarrow \infty} \sum_{m=1}^M \frac{x_m}{M}. \quad (38)$$

As we obviously do not have infinite time nor infinite quantum computers (which could be run in parallel), we must truncate our number of measurement M to a finite, but sufficiently large number. More precisely, for precision ϵ , each expectation estimation subroutine within VQE requires $\mathcal{O}(1/\epsilon^2)$ samples from circuits with depth $\mathcal{O}(1)$.

Quantum computing and solving the eigenvalue problem for the Lipkin model

We turn now to a simpler variant of the Lipkin model without the W -term and a total spin of $J = 1$ only as maximum value of the spin. This corresponds to a system with $N = 2$ particles (fermions in our case). Our Hamiltonian is given by the quasispin operators (see below)

$$\hat{H} = \epsilon \hat{J}_z - \frac{1}{2} V (\hat{J}_+ \hat{J}_+ + \hat{J}_- \hat{J}_-).$$

As discussed previously, the quasispin operators act like lowering and raising angular momentum operators.

With these properties we can calculate the Hamiltonian matrix for the Lipkin model by computing the various matrix elements

$$\langle J J_z | H | J J'_z \rangle, \quad (39)$$

where the non-zero elements are given by

$$\begin{aligned} \langle J J_z | H | J J'_z \rangle &= \epsilon J_z \\ \langle J J_z | H | J J'_z \pm 2 \rangle &= \langle J J_z \pm 2 | H | J J'_z \rangle \\ &= -\frac{1}{2} V C, \end{aligned}$$

where C is the Clebsch-Gordan coefficients (from the raising and lowering operators) one gets when J_\pm^2 operates on the state $|J J_z\rangle$. Using the above definitions we can calculate the exact solution to the Lipkin model.

With the V -interaction terms, we obtain the following Hamiltonian matrix

$$\begin{pmatrix} -\epsilon & 0 & -V \\ 0 & 0 & 0 \\ -V & 0 & \epsilon \end{pmatrix} \quad (40)$$

The following **python** code sets up the above matrix and finds the pertinent eigenvalues.

```
import numpy as np
import qiskit
from qiskit.visualization import circuit_drawer
from qiskit.quantum_info import Statevector
from matplotlib.pyplot import figure
from qiskit import QuantumRegister, QuantumCircuit, ClassicalRegister, Aer, assemble
from qiskit.providers.aer.noise import NoiseModel
import pylatexenc
from qiskit.algorithms import VQE
from qiskit.utils import QuantumInstance
from qiskit.opflow import X, Z, I, Y
from qiskit.circuit import Parameter
from qiskit.algorithms.optimizers import ADAM
from qiskit.opflow import AerPauliExpectation
from qiskit import IBMQ
import cmath
import pandas as pd
from scipy.sparse import diags
import numpy.linalg as LA
import matplotlib.pyplot as plt
from IPython.display import Image
import warnings
warnings.filterwarnings('ignore')
pi=np.pi

#function that sorts eigenvalues with its eigenvectors in accending order
def eigen(A):
    eigenValues, eigenVectors = LA.eig(A)
    idx = np.argsort(eigenValues)
    eigenValues = eigenValues[idx]
    eigenVectors = eigenVectors[:,idx]
    return (eigenValues, eigenVectors)

#one body expectation value
def one_body(E,N):
    k = N/2
    m = np.arange(-k,k+1,1) # Since the collective space is Omega+1
    return E*np.diag(m) #return a matrix where its diagonal elemens are epsilon*K_0

#two body expectation value
def two_body(V,N):
    k = N/2
    m = np.arange(-k,k+1,1)
    left =np.zeros(len(m)-2,dtype=complex)
    right = np.zeros(len(m)-2,dtype=complex)
    diag = np.zeros(len(m),dtype=complex)
    for i in range(len(left)):
        CG = cmath.sqrt(k*(k+1)-(m[i]+2)*(m[i]+1))*cmath.sqrt(k*(k+1)-m[i]*(m[i]+1)) #calculate Clebs
        left[i] = CG
```

```

        right[i] = CG
    k = [left,diag,right]
    offset = [-2,0,2]
    return -0.5*V*diags(k,offset).toarray() #return a matrix where its off diagonal elements are (1/

#full expectation value
def quasi_spin(E,V,N):
    ob = one_body(E,N)
    tb = two_body(V,N)
    H = ob+tb
    e,v = eigen(H) # find the eigenvalues of the Hamiltonian
    return e,H

#converts chi to V
def Vp(E,omega,chi):
    return (chi*E)/(omega-1)

#parameters
E = 1
chi = np.arange(0,2.1,0.1)
omega = 2

EV0 = []
EV1 = []
EV2 = []
Ham = []
for i in chi:
    v = Vp(E,omega,i)
    EigenV,H = quasi_spin(E,v,omega) #return eigenvalues and Hamiltonian
    Ham.append(H)
    EV0.append(EigenV[0])
    EV1.append(EigenV[1])
    EV2.append(EigenV[2])

matrix = pd.DataFrame(Ham[5].real)
print('Hamiltonian matrix')
matrix.head()

plt.plot(chi,EV0)
plt.plot(chi,EV1)
plt.plot(chi,EV2)
plt.xlabel('$\chi$')
plt.ylabel('Energy')
plt.title('$\Omega=2$ exact Lipkin Model')

```

Quantum Circuit, rewriting the Lipkin model in terms of Pauli matrices. To solve the Lipkin model on a quantum computer we have to solve Schrodinger's equation. To achieve this, we will use the Variational Quantum Eigensolver (VQE) [discussed earlier](#).

Before we proceed however, we need to rewrite the quasispin operators in terms of Pauli spin matrices/operators.

We take the liberty here of reminding you of some of the derivations done previously. We defined the number operator as

$$N = \sum_{n\sigma} a_{n\sigma}^\dagger a_{n\sigma},$$

which commutes with the Lipkin Hamiltonian. This can be seen by examining the Lipkin model Hamiltonian and noticing that the one-body part simply counts particles while the two-body term moves particles in pairs. Thus, the Hamiltonian conserves particle number. To find more symmetries we rewrote the Lipkin Hamiltonian in terms of $SU(2)$ quasispin operators

$$H = \epsilon J_z + \frac{1}{2} V (J_+^2 + J_-^2), \quad (41)$$

via the mappings

$$J_z = \sum_n j_z^{(n)},$$

and

$$J_\pm = \sum_n j_\pm^{(n)},$$

where we have the onebody operators

$$j_z^{(n)} = \frac{1}{2} \sum_\sigma \sigma a_{n\sigma}^\dagger a_{n\sigma},$$

and

$$j_\pm^{(n)} = a_{n\pm}^\dagger a_{n\mp}.$$

These operators obey the $SU(2)$ commutation relations

$$[J_+, J_-] = 2J_z,$$

and

$$[J_z, J_\pm] = \pm J_\pm.$$

Here the ladder operators are defined as $J_\pm = J_x \pm iJ_y$. With this rewriting, we can see that the total spin operator J^2 , which is defined as

$$J^2 = J_x^2 + J_y^2 + J_z^2 = \frac{1}{2} \{J_+, J_-\} + J_z^2,$$

commutes with the Hamiltonian since the Hamiltonian. We note also that the rotation operator

$$R = e^{i\phi J_z},$$

commutes with the Hamiltonian, which can be explained as follows. Writing J_z as

$$J_z = \frac{1}{2} (N_+ - N_-),$$

where $N_{\pm} = \sum_{n\pm} a_{n\pm}^{\dagger} a_{n\pm}$, allows us to see that it measures half the difference between the number of particles in the upper and lower levels. Thus, the possible eigenvalues r of the signature operator are

$$r = +1, j_z = 2n \quad (42)$$

$$r = +i, j_z = 2n + \frac{1}{2} \quad (43)$$

$$r = -1, j_z = 2n + 1 \quad (44)$$

$$r = -i, j_z = 2n + \frac{3}{2} \quad (45)$$

$$(46)$$

for $n \in \mathbb{Z}$. Note that r is real or imaginary if the number of particles N is even or odd, respectively. Since, as discussed above, the Lipkin Hamiltonian conserves N , r cannot jump between being real and imaginary. Additionally, because particles must be moved in pairs, and J_z measures half the difference between particles in the upper and lower levels, j_z can only change by as

$$j_z \rightarrow \frac{1}{2}[(N_+ \pm 2n) - (N_- \mp 2n)]$$

or $j_z \rightarrow J_z \pm 2n$.

To solve the Lipkin model with a quantum computer, the first step is to map the system to a set of qubits. We will restrict ourselves here to the half-filled case where the number of particles N equals the degeneracy of the states Ω . One could assign each possible state (n, σ) a qubit such that the qubit being in the state $|1\rangle$ or $|0\rangle$ would imply that the state (n, σ) is occupied or unoccupied, respectively. This mapping scheme (which we will call occupation mapping) requires 2Ω qubits.

However, because there are only two energy levels in the Lipkin model, any other natural mapping is possible. In this mapping scheme (which we will call level mapping) each doublet $((n, +1), (n, -1))$ would be assigned a qubit such that the qubit being in the state $|0\rangle$ or $|1\rangle$ would imply that the particle is in the $(n, +1)$ or $(n, -1)$ state, respectively. Note that these are the only two possible configurations of the doublet as we are restricting ourselves to the half-filled case and the Lipkin Hamiltonian only moves particles between energy levels, not degenerate states. Thus the level mapping only requires Ω qubits which is half that of the occupation mapping. Additionally, any ansatz that would restrict the minimization search to the correct subspace of constant Hamming weight N requires at most, only two qubit gates. This is because moving a pair of particles in this scheme only changes the state of two doublets (and therefore qubits). That is, it only takes a two-qubit gate to change between the states $|00\rangle$ and $|11\rangle$, for example. As an efficient decomposition two-qubit gates is known, the ansatz for this mapping would be shorter (and thus less noisy) than that of the previous mapping.

One could imagine a third mapping scheme which would require even less qubits in which each of the possible states in the spin basis $|JJ_z\rangle$ is mapped to

a single qubit. In this spin mapping, there are only $2J + 1$ possible states (since $J_z = -J, -J + 1, \dots, J - 1, J$) for each value of J . And, since the Hamiltonian is block diagonal (with a different block for each J) the eigenvalues of the Hamiltonian are simply the eigenvalues of each block, which may be calculated separately. Since the maximum value of J is $J_{\max} = N/2$, the largest number of qubits would be $2J_{\max} + 1 = N + 1$. However, $\lfloor N/2 \rfloor$ different circuit would need to be used for minimization for all possible values of J , to explore the entire Hilbert space. (The minimum of the set of minimum energies that each circuit finds would be the ground state energy of the entire system.) This increases, linearly, the amount of time required to find the ground state energy.

After reviewing the three possible mappings, it is our view that the level mapping is the best suited for NISQ era devices given its low qubit count and ability to search the entire relevant Hilbert space with one circuit (which reduces time to solution) and the fact that at most, only two-qubit gates are required of the ansatz, leading to shorter depth (and thus less noisy) circuits.

The Hamiltonian takes the form

$$H = \epsilon J_z + \frac{1}{2} V (J_+^2 + J_-^2). \quad (47)$$

Plugging the mapping from the total J operators to the individual one-body j operators yields

$$H = \epsilon \sum_n j_z^{(n)} + \frac{1}{2} V \left[\left(\sum_n j_+^{(n)} \right)^2 + \left(\sum_n j_-^{(n)} \right)^2 \right] \quad (48)$$

$$= \epsilon \sum_n j_z^{(n)} + \frac{1}{2} V \sum_{n,m} \left(j_+^{(n)} j_+^{(m)} + j_-^{(n)} j_-^{(m)} \right) \quad (49)$$

$$= \epsilon \sum_n j_z^{(n)} + 2V \sum_{n < m} \left(j_x^{(n)} j_x^{(m)} - j_y^{(n)} j_y^{(m)} \right), \quad (50)$$

where we have used the definitions

$$j_{\pm}^{(n)} = j_x^{(n)} \pm i j_y^{(n)}.$$

To convert to Pauli matrices, we make the transformations

$$j_x^{(n)} \rightarrow X_n/2,$$

and

$$j_y^{(n)} \rightarrow Y_n/2,$$

and finally

$$j_z^{(n)} \rightarrow Z_n/2,$$

which preserve the above $SU(2)$ commutation relations. The factor of $1/2$ is due to the eigenvalues of the Pauli matrices being ± 1 while we are dealing with spin $1/2$ particles.

This transforms our Hamiltonian into

$$H = \frac{1}{2}\epsilon \sum_{k=1}^n Z_k + \frac{1}{2}V \sum_{n \neq j=1}^N (X_k X_j - Y_k Y_j).$$

With this form, we can clearly see that the first (one-body) term in the Hamiltonian returns the energy $-\epsilon/2$ or $+\epsilon/2$ if the qubit representing the particle of a doublet is in the ground ($|1\rangle$) or excited ($|0\rangle$) state, respectively. The action of the second (two-body) term in the Hamiltonian can be determined by noting that

$$\frac{1}{2}(XX - YY)|00\rangle = |11\rangle, \quad (51)$$

$$\frac{1}{2}(XX - YY)|01\rangle = 0, \quad (52)$$

$$\frac{1}{2}(XX - YY)|10\rangle = 0, \quad (53)$$

$$\frac{1}{2}(XX - YY)|11\rangle = |00\rangle. \quad (54)$$

That is, the two-body term moves a pair of particles between the ground states $|00\rangle$ and the excited states $|11\rangle$ of their respective doublets.

To construct an efficient ansatz, we must determine the subspace within which the Hamiltonian lives. To begin, note that particles are only ever moved between energy levels in pairs. Further, note that the Hamiltonian's coefficients (ϵ and V) are state independent (do not depend on the indices n or m) as the states labeled by these indices are degenerate and thus have the same energy level. Thus, the Hamiltonian treats all states with the same number of excited particles (Hamming weight of the state) as the same. Therefore, the following ansatz forms exactly cover the subspace within which the N -degenerate Hamiltonian explores:

$$|\psi_{\text{even}}\rangle = \sum_{k=0}^{\lfloor n/2 \rfloor} c_{2k} |D_{2k}^n\rangle, \quad (55)$$

$$|\psi_{\text{odd}}\rangle = \sum_{k=0}^{\lfloor n/2 \rfloor} c_{2k+1} |D_{2k+1}^n\rangle. \quad (56)$$

Here $|D_k^n\rangle$ represents a Dicke state which is defined as equal superposition of all n -qubit states with Hamming weight k . That is

$$|D_k^n\rangle = \frac{1}{\sqrt{\binom{n}{k}}} \sum_{x \in h_k^n} |x\rangle, \quad (57)$$

where $h_k^n = \{|x\rangle \mid l(x) = n, \text{wt}(x) = k\}$. There are two ways we can think of to prepare such ansatz: The first is to prepare them exactly as it is known

how to deterministically prepare Dicke states with linear depth. The reference provides an algorithm for preparing a set of gates U_k^n that prepares a Dicke state from a product state of Hamming weight k ; that is

$$U_k^n |1\rangle^{\otimes k} |0\rangle^{\otimes n-k} = |D_k^n\rangle. \quad (58)$$

It then describes how to one can create an arbitrary superposition of Dicke states, which we modify here to restrict ourselves to a Hamming weight of constant parity. The circuit to construct such a state (for the $k = 6$ case, as an example) is discussed next week.

Implementing the VQE method, one qubit system

We start with a simple 2×2 Hamiltonian matrix expressed in terms of Pauli X and Z matrices, as discussed in the project text.

We define a symmetric matrix $H \in \mathbb{R}^{2 \times 2}$

$$H = \begin{bmatrix} H_{11} & H_{12} \\ H_{21} & H_{22} \end{bmatrix},$$

We let $H = H_0 + H_I$, where

$$H_0 = \begin{bmatrix} E_1 & 0 \\ 0 & E_2 \end{bmatrix},$$

is a diagonal matrix. Similarly,

$$H_I = \begin{bmatrix} V_{11} & V_{12} \\ V_{21} & V_{22} \end{bmatrix},$$

where V_{ij} represent various interaction matrix elements. We can view H_0 as the non-interacting solution

$$H_0 |0\rangle = E_1 |0\rangle, \quad (59)$$

and

$$H_0 |1\rangle = E_2 |1\rangle, \quad (60)$$

where we have defined the orthogonal computational one-qubit basis states $|0\rangle$ and $|1\rangle$.

We rewrite H (and H_0 and H_I) via Pauli matrices

$$H_0 = \mathcal{E}I + \Omega\sigma_z, \quad \mathcal{E} = \frac{E_1 + E_2}{2}, \quad \Omega = \frac{E_1 - E_2}{2},$$

and

$$H_I = cI + \omega_z\sigma_z + \omega_x\sigma_x,$$

with $c = (V_{11} + V_{22})/2$, $\omega_z = (V_{11} - V_{22})/2$ and $\omega_x = V_{12} = V_{21}$. We let our Hamiltonian depend linearly on a strength parameter λ

$$H = H_0 + \lambda H_I,$$

with $\lambda \in [0, 1]$, where the limits $\lambda = 0$ and $\lambda = 1$ represent the non-interacting (or unperturbed) and fully interacting system, respectively. The model is an eigenvalue problem with only two available states.

Here we set the parameters $E_1 = 0$, $E_2 = 4$, $V_{11} = -V_{22} = 3$ and $V_{12} = V_{21} = 0.2$.

The non-interacting solutions represent our computational basis. Pertinent to our choice of parameters, is that at $\lambda \geq 2/3$, the lowest eigenstate is dominated by $|1\rangle$ while the upper is $|0\rangle$. At $\lambda = 1$ the $|0\rangle$ mixing of the lowest eigenvalue is 1% while for $\lambda \leq 2/3$ we have a $|0\rangle$ component of more than 90%. The character of the eigenvectors has therefore been interchanged when passing $z = 2/3$. The value of the parameter V_{12} represents the strength of the coupling between the two states.

Setting up the matrix

```
from matplotlib import pyplot as plt
import numpy as np
dim = 2
Hamiltonian = np.zeros((dim,dim))
e0 = 0.0
e1 = 4.0
Xnondiag = 0.20
Xdiag = 3.0
Eigenvalue = np.zeros(dim)
# setting up the Hamiltonian
Hamiltonian[0,0] = Xdiag+e0
Hamiltonian[0,1] = Xnondiag
Hamiltonian[1,0] = Hamiltonian[0,1]
Hamiltonian[1,1] = e1-Xdiag
# diagonalize and obtain eigenvalues, not necessarily sorted
EigValues, EigVectors = np.linalg.eig(Hamiltonian)
permute = EigValues.argsort()
EigValues = EigValues[permute]
# print only the lowest eigenvalue
print(EigValues[0])
```

Now rewrite it in terms of the identity matrix and the Pauli matrix X and Z

```
# Now rewrite it in terms of the identity matrix and the Pauli matrix X and Z
X = np.array([[0,1],[1,0]])
Y = np.array([[0,-1j],[1j,0]])
Z = np.array([[1,0],[0,-1]])
# identity matrix
I = np.array([[1,0],[0,1]])

epsilon = (e0+e1)*0.5; omega = (e0-e1)*0.5
c = 0.0; omega_z=Xdiag; omega_x = Xnondiag
Hamiltonian = (epsilon+c)*I+(omega_z+omega)*Z+omega_x*X
EigValues, EigVectors = np.linalg.eig(Hamiltonian)
permute = EigValues.argsort()
EigValues = EigValues[permute]
# print only the lowest eigenvalue
print(EigValues[0])
```

Implementing the VQE

For a one-qubit system we can reach every point on the Bloch sphere (as discussed earlier) with a rotation about the x -axis and the y -axis.

We can express this mathematically through the following operations (see whiteboard for the drawing), giving us a new state $|\psi\rangle$

$$|\psi\rangle = R_y(\phi)R_x(\theta)|0\rangle.$$

We can produce multiple ansatzes for the new state in terms of the angles θ and ϕ . With these ansatzes we can in turn calculate the expectation value of the above Hamiltonian, now rewritten in terms of various Pauli matrices (and thereby gates), that is compute

$$\langle\psi|(c + \mathcal{E})\mathbf{I} + (\Omega + \omega_z)\boldsymbol{\sigma}_z + \omega_x\boldsymbol{\sigma}_x|\psi\rangle.$$

We can now set up a series of ansatzes for $|\psi\rangle$ as function of the angles θ and ϕ and find thereafter the variational minimum using for example a gradient descent method.

To do so, we need to remind ourselves about the mathematical expressions for the rotational matrices/operators.

$$R_x(\theta) = \cos \frac{\theta}{2} \mathbf{I} - \imath \sin \frac{\theta}{2} \boldsymbol{\sigma}_x,$$

and

$$R_y(\phi) = \cos \frac{\phi}{2} \mathbf{I} - \imath \sin \frac{\phi}{2} \boldsymbol{\sigma}_y.$$

```
# define the rotation matrices
# Define angles theta and phi
theta = 0.5*np.pi; phi = 0.2*np.pi
Rx = np.cos(theta*0.5)*I-1j*np.sin(theta*0.5)*X
Ry = np.cos(phi*0.5)*I-1j*np.sin(phi*0.5)*Y
#define basis states
basis0 = np.array([1,0])
basis1 = np.array([0,1])

NewBasis = Ry @ Rx @ basis0
print(NewBasis)
# Compute the expectation value
#Note hermitian conjugation
Energy = NewBasis.conj().T @ Hamiltonian @ NewBasis
print(Energy)
```

Not an impressive results. We set up now a loop over many angles θ and ϕ and compute the energies

```
# define a number of angles
n = 20
angle = np.arange(0,180,10)
n = np.size(angle)
```

```

ExpectationValues = np.zeros((n,n))
for i in range (n):
    theta = np.pi*angle[i]/180.0
    Rx = np.cos(theta*0.5)*I-1j*np.sin(theta*0.5)*X
    for j in range (n):
        phi = np.pi*angle[j]/180.0
        Ry = np.cos(phi*0.5)*I-1j*np.sin(phi*0.5)*Y
        NewBasis = Ry @ Rx @ basis0
        Energy = NewBasis.conj().T @ Hamiltonian @ NewBasis
        Edifference=abs(np.real(EigValues[0]-Energy))
        ExpectationValues[i,j]=Edifference

print(np.min(ExpectationValues))

```

Clearly, this is not the very best way of proceeding. Rather, here we would compute the gradient and thereby find the minimum as function of the angles θ and ϕ . Furthermore, in setting up the angles, a better practice is to select random values for these.

For the lectures of April 17-21, we will add code example using gradient descent for the one- and two-qubit case. We will follow <https://journals.aps.org/pr/abstract/10.1103/PhysRevA.99.032331> as a guideline to calculate gradients of the Hamiltonian.

A smarter way of doing this

The above approach means that we are setting up several matrix-matrix and matrix-vector multiplications. Although straight forward it is not the most efficient way of doing this, in particular in case the matrices become large (and sparse). But there are some more important issues.

In a physical realization of these systems we cannot just multiply the state with the Hamiltonian. When performing a measurement we can only measure in one particular direction. For the computational basis states which we have, $|0\rangle$ and $|1\rangle$, we have to measure along the bases of the Pauli matrices and reconstruct the eigenvalues from these measurements.

From our earlier discussions we know that the Pauli Z matrix has the above basis states as eigen states through

$$\sigma_z|0\rangle = Z|0\rangle = +1|0\rangle,$$

and

$$\sigma_z|1\rangle = Z|1\rangle = -1|1\rangle,$$

with eigenvalue -1 .

For the Pauli X matrix on the other hand we have

$$\sigma_x|0\rangle = X|0\rangle = +1|1\rangle,$$

and

$$\sigma_x|1\rangle = X|1\rangle = -1|0\rangle,$$

with eigenvalues 1 in both cases. The latter two equations tell us that the computational basis we have chosen, and in which we will prepare our states, is not an eigenbasis of the σ_x matrix.

We will thus try to rewrite the Pauli X matrix in terms of a Pauli Z matrix. Fortunately this can be done using the Hadamard matrix twice, that is

$$\mathbf{X} = \sigma_x = \mathbf{H}\mathbf{Z}\mathbf{H}.$$

The Pauli Y matrix can be written as

$$\mathbf{Y} = \sigma_y = \mathbf{H}\mathbf{S}^\dagger\mathbf{Z}\mathbf{H}\mathbf{S},$$

where \mathbf{S} is the phase matrix

$$\mathbf{S} = \begin{bmatrix} 1 & 0 \\ 0 & i \end{bmatrix}.$$

From here and on we will denote the Pauli matrices by X , Y and Z and we can write the expectation value of the Hamiltonian as

$$\langle\psi|(c + \mathcal{E})\mathbf{I} + (\Omega + \omega_z)\mathbf{Z} + \omega_x\mathbf{H}\mathbf{Z}\mathbf{H}|\psi\rangle,$$

which we can rewrite as

$$(c + \mathcal{E})\langle\psi|\mathbf{I}|\psi\rangle + (\Omega + \omega_z)\langle\psi|\mathbf{Z}|\psi\rangle + \omega_x\langle\psi\mathbf{H}|\mathbf{Z}|\mathbf{H}\psi\rangle.$$

The first and second term are too easy to perform a measurement on since we just need to compute $\langle\psi|\mathbf{I}|\psi\rangle$ and $\langle\psi|\mathbf{Z}|\psi\rangle$. For the final term we need just to add the action of the Hadamard matrix and we are done.

Two-qubit system and the VQE

We extend now the system to a two-qubit system with the following computational basis states and Hamiltonian matrix written out in terms of Pauli spin matrices.

This system can be thought of as composed of two subsystems A and B . Each subsystem has computational basis states

$$|0\rangle_{A,B} = \begin{bmatrix} 1 & 0 \end{bmatrix}^T \quad |1\rangle_{A,B} = \begin{bmatrix} 0 & 1 \end{bmatrix}^T.$$

The subsystems could represent single particles or composite many-particle systems of a given symmetry. This leads to the many-body computational basis states

$$|00\rangle = |0\rangle_A \otimes |0\rangle_B = \begin{bmatrix} 1 & 0 & 0 & 0 \end{bmatrix}^T,$$

and

$$|01\rangle = |0\rangle_A \otimes |1\rangle_B = \begin{bmatrix} 0 & 1 & 0 & 0 \end{bmatrix}^T,$$

and

$$|10\rangle = |1\rangle_A \otimes |0\rangle_B = \begin{bmatrix} 0 & 0 & 1 & 0 \end{bmatrix}^T,$$

and finally

$$|11\rangle = |1\rangle_A \otimes |1\rangle_B = \begin{bmatrix} 0 & 0 & 0 & 1 \end{bmatrix}^T.$$

These computational basis states define also the eigenstates of the non-interacting Hamiltonian

$$H_0|00\rangle = \epsilon_{00}|00\rangle,$$

$$H_0|10\rangle = \epsilon_{10}|10\rangle,$$

$$H_0|01\rangle = \epsilon_{01}|01\rangle,$$

and

$$H_0|11\rangle = \epsilon_{11}|11\rangle.$$

The interacting part of the Hamiltonian H_I is given by the tensor product of two σ_x and σ_z matrices, respectively, that is

$$H_I = H_x \sigma_x \otimes \sigma_x + H_z \sigma_z \otimes \sigma_z,$$

where H_x and H_z are interaction strength parameters. Our final Hamiltonian matrix is given by

$$\mathbf{H} = \begin{bmatrix} \epsilon_{00} + H_z & 0 & 0 & H_x \\ 0 & \epsilon_{10} - H_z & H_x & 0 \\ 0 & H_x & \epsilon_{01} - H_z & 0 \\ H_x & 0 & 0 & \epsilon_{11} + H_z \end{bmatrix}.$$

The four eigenstates of the above Hamiltonian matrix can in turn be used to define density matrices. As an example, the density matrix of the first eigenstate (lowest energy E_0) Ψ_0 is

$$\rho_0 = (\alpha_{00}|00\rangle\langle 00| + \alpha_{10}|10\rangle\langle 10| + \alpha_{01}|01\rangle\langle 01| + \alpha_{11}|11\rangle\langle 11|),$$

where the coefficients α_{ij} are the eigenvector coefficients resulting from the solution of the above eigenvalue problem.

Switching to Qiskit (notes developed by Stian Bilek)

```
import numpy as np
import qiskit as qk
from scipy.optimize import minimize

# Initialize registers and circuit
n_qubits = 1 #Number of qubits
n_cbits = 1 #Number of classical bits (the number of qubits you want to measure at the end of the
qreg = qk.QuantumRegister(n_qubits) #Create a quantum register
creg = qk.ClassicalRegister(n_cbits) #Create a classical register
circuit = qk.QuantumCircuit(qreg,creg) #Create your quantum circuit
circuit.draw() #Draw circuit. It is empty
# Perform operations on qubit
circuit.x(qreg[0]) #Applies a Pauli X gate to the first qubit in the quantum register
circuit.draw()
# Chose a qubit to measure and encode the results to a classical bit
```

```

#Measure the first qubit in the quantum register
#and encode the results to the first qubit in the classical register
circuit.measure(qreg[0],creg[0])
circuit.draw()
# Execute circuit

backend = qk.Aer.get_backend('qasm_simulator')
#This is the device you want to use. It is an ideal simulation of a quantum device

job = backend.run(circuit,shots=1000) #Run the circuit 1000 times
result = job.result()
counts = result.get_counts()
print(counts)
circuit.clear()
circuit.draw()

circuit.h(qreg[0]) #Apply a Hadamard gate to the first qubit of the quantum register
circuit.measure(qreg,creg)
print(circuit.draw())
job = backend.run(circuit,shots=1000)
result = job.result()
counts = result.get_counts()
print(counts)
circuit.clear()

```

Create a two-qubit circuit and set up a Bell state

```

n_qubits = 2
n_cbits = 2
qreg = qk.QuantumRegister(n_qubits)
creg = qk.ClassicalRegister(n_cbits)
circuit = qk.QuantumCircuit(qreg,creg)
circuit.draw()
circuit.h(qreg[0])
circuit.cx(qreg[0],qreg[1])

```

This is a controlled operation. Apply a Pauli X gate to the second qubit ($qreg[1]$) if the first qubit ($qreg[0]$) is in the $|1\rangle$ state. Else do nothing

```

circuit.draw()
circuit.measure(qreg,creg)
circuit.draw()
job = backend.run(circuit,shots=1000)
result = job.result()
counts = result.get_counts()
print(counts)
circuit.clear()

```

Apply rotation to qubit

```

theta = np.pi/3
circuit.rx(theta, qreg[0]) #R_x(theta) rotation on the first qubit (qreg[0])
circuit.measure(qreg,creg)
print(circuit.draw())
job = backend.run(circuit,shots=1000)

```

```

result = job.result()
counts = result.get_counts()
circuit.clear()
print(counts)

```

Find the lowest eigenvalue of

$$H = c_1 Z_0 + c_2 Z_1 + c_3 X_0 Y_1$$

We will use

$$\langle \psi | H | \psi \rangle = c_1 \langle \psi | Z_0 | \psi \rangle + c_2 \langle \psi | Z_1 | \psi \rangle + c_3 \langle \psi | X_0 Y_1 | \psi \rangle$$

```

I = np.eye(2)
X = np.array([[0,1],[1,0]])
Y = np.array([[0,-1j],[1j,0]])
Z = np.array([[1,0],[0,-1]])
H = np.kron(Z,I) + np.kron(I,Z) + np.kron(X,Y)
eigvals,eigvecs = np.linalg.eigh(H)
print(eigvals[0])

c_1 = 1
c_2 = 1
c_3 = 1

h_1 = [c_1,[0],['z']]
h_2 = [c_2,[1],['z']]
h_3 = [c_3,[0,1],['x','x']]
H = [h_1,h_2,h_3]

```

Create ansatz

```

def ansatz(theta,n_qubits):
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)
    for i in range(n_qubits):
        circuit.ry(theta[i],qreg[i])
    for i in range(n_qubits-1):
        circuit.cx(qreg[i],qreg[i+1])
    return(circuit)
qreg = qk.QuantumRegister(n_qubits)
circuit = qk.QuantumCircuit(qreg)
circuit.h(qreg[:2])
print('Before ansatz')
print(circuit.draw())
theta = np.random.randn(2)
n_qubits = 2
circuit = circuit.compose(ansatz(theta,n_qubits))
print('After ansatz')
circuit.draw()

```

Change measurement basis

```

def basis_change(h_i,n_qubits):
    qreg = qk.QuantumRegister(n_qubits)

```

```

circuit = qk.QuantumCircuit(qreg)

for qubit, operator in zip(h_i[1], h_i[2]):
    if operator == 'x':
        circuit.h(qreg[qubit])
    if operator == 'y':
        circuit.sdg(qreg[qubit])
        circuit.h(qreg[qubit])
    return(circuit)
n_qubits = 2
qreg = qk.QuantumRegister(n_qubits)
circuit = qk.QuantumCircuit(qreg)
theta = np.random.randn(n_qubits)
circuit = circuit.compose(ansatz(theta, n_qubits))
print('Ansatz circuit')
circuit.draw()
circuit = circuit.compose(basis_change(H[2], n_qubits))
print('After basis transformation:')
print(circuit.draw())

```

Get energy for given rotational parameters, theta

```

def get_energy(theta):
    n_qubits = 2
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)
    circuit = circuit.compose(ansatz(theta, n_qubits))
    circuit_list = []
    for idx, h_i in enumerate(H):
        basis_change_circuit = basis_change(h_i, n_qubits)
        new_circuit = circuit.compose(basis_change_circuit)
        creg = qk.ClassicalRegister(len(h_i[1]))
        new_circuit.add_register(creg)
        new_circuit.measure(qreg[h_i[1]], creg)
        circuit_list.append(new_circuit)
    shots = 10000
    job = backend.run(circuit_list, shots=shots)
    E = np.zeros(len(circuit_list))
    for i in range(len(circuit_list)):
        result = job.result()
        counts = result.get_counts(i)
        for key, value in counts.items():
            e = 1
            for bit in key:
                if bit == '0':
                    e *= 1
                if bit == '1':
                    e *= -1
            E[i] += e*value
        E[i] *= H[i][0]
    E /= shots
    return(np.sum(E))

theta = np.random.randn(2)
get_energy(theta)

```

Minimize energy with Scipy

```

theta = np.random.randn(2)
res = minimize(get_energy, theta, method='Powell',tol=1e-12)
get_energy(res.x)

### We might need a more flexible ansatz

def ansatz(theta,n_qubits):
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)
    idx = 0
    for i in range(n_qubits):
        circuit.ry(theta[idx],qreg[i])
        idx += 1
    for i in range(n_qubits-1):
        circuit.cx(qreg[i],qreg[i+1])
    for i in range(n_qubits):
        circuit.rx(theta[idx],qreg[i])
        idx += 1
    for i in range(n_qubits-1):
        circuit.cx(qreg[i],qreg[i+1])
    return(circuit)
theta = np.random.randn(4)
res = minimize(get_energy, theta, method='Powell',tol=1e-16)
get_energy(res.x)

```

Then we minimize energy with gradient descent

$$\frac{\partial E(\theta_1, \dots, \theta_i, \dots, \theta_p)}{\partial \theta_i} = \frac{E(\theta_1, \dots, \theta_i + \pi/2, \dots, \theta_p) - E(\theta_1, \dots, \theta_i - \pi/2, \dots, \theta_p)}{2}.$$

```

epochs = 200
theta = np.random.randn(4)
for epoch in range(epochs):
    print(epoch,get_energy(theta))
    grad = np.zeros_like(theta)
    for idx in range(theta.shape[0]):
        theta_temp = theta.copy()
        theta_temp[idx] += np.pi/2
        E_plus = get_energy(theta_temp)
        theta_temp[idx] -= np.pi
        E_minus = get_energy(theta_temp)
        grad[idx] = (E_plus - E_minus)/2
    theta -= 0.1*grad

```

VQE and efficient computations of gradients

We start with a reminder on the VQE method with applications to the one-qubit system. We discussed this to some detail during the week of March 27-31. Here we revisit the one-qubit system and develop a VQE code for studying this system using gradient descent as a method to optimize the variational ansatz.

We start with a simple 2×2 Hamiltonian matrix expressed in terms of Pauli X and Z matrices, as discussed in the project text.

We define a symmetric matrix $H \in \mathbb{R}^{2 \times 2}$

$$H = \begin{bmatrix} H_{11} & H_{12} \\ H_{21} & H_{22} \end{bmatrix},$$

We let $H = H_0 + H_I$, where

$$H_0 = \begin{bmatrix} E_1 & 0 \\ 0 & E_2 \end{bmatrix},$$

is a diagonal matrix. Similarly,

$$H_I = \begin{bmatrix} V_{11} & V_{12} \\ V_{21} & V_{22} \end{bmatrix},$$

where V_{ij} represent various interaction matrix elements. We can view H_0 as the non-interacting solution

$$H_0|0\rangle = E_1|0\rangle, \quad (61)$$

and

$$H_0|1\rangle = E_2|1\rangle, \quad (62)$$

where we have defined the orthogonal computational one-qubit basis states $|0\rangle$ and $|1\rangle$.

We rewrite H (and H_0 and H_I) via Pauli matrices

$$H_0 = \mathcal{E}I + \Omega\sigma_z, \quad \mathcal{E} = \frac{E_1 + E_2}{2}, \quad \Omega = \frac{E_1 - E_2}{2},$$

and

$$H_I = cI + \omega_z\sigma_z + \omega_x\sigma_x,$$

with $c = (V_{11} + V_{22})/2$, $\omega_z = (V_{11} - V_{22})/2$ and $\omega_x = V_{12} = V_{21}$. We let our Hamiltonian depend linearly on a strength parameter λ

$$H = H_0 + \lambda H_I,$$

with $\lambda \in [0, 1]$, where the limits $\lambda = 0$ and $\lambda = 1$ represent the non-interacting (or unperturbed) and fully interacting system, respectively. The model is an eigenvalue problem with only two available states.

Here we set the parameters $E_1 = 0$, $E_2 = 4$, $V_{11} = -V_{22} = 3$ and $V_{12} = V_{21} = 0.2$.

The non-interacting solutions represent our computational basis. Pertinent to our choice of parameters, is that at $\lambda \geq 2/3$, the lowest eigenstate is dominated by $|1\rangle$ while the upper is $|0\rangle$. At $\lambda = 1$ the $|0\rangle$ mixing of the lowest eigenvalue is 1% while for $\lambda \leq 2/3$ we have a $|0\rangle$ component of more than 90%. The character of the eigenvectors has therefore been interchanged when passing $\lambda = 2/3$. The value of the parameter V_{12} represents the strength of the coupling between the two states.

Setting up the matrix

```
from matplotlib import pyplot as plt
import numpy as np
dim = 2
Hamiltonian = np.zeros((dim,dim))
e0 = 0.0
e1 = 4.0
Xnondiag = 0.20
Xdiag = 3.0
Eigenvalue = np.zeros(dim)
# setting up the Hamiltonian
Hamiltonian[0,0] = Xdiag+e0
Hamiltonian[0,1] = Xnondiag
Hamiltonian[1,0] = Hamiltonian[0,1]
Hamiltonian[1,1] = e1-Xdiag
# diagonalize and obtain eigenvalues, not necessarily sorted
EigValues, EigVectors = np.linalg.eig(Hamiltonian)
permute = EigValues.argsort()
EigValues = EigValues[permute]
# print only the lowest eigenvalue
print(EigValues[0])
```

Now rewrite it in terms of the identity matrix and the Pauli matrix X and Z

```
# Now rewrite it in terms of the identity matrix and the Pauli matrix X and Z
X = np.array([[0,1],[1,0]])
Y = np.array([[0,-1j],[1j,0]])
Z = np.array([[1,0],[0,-1]])
# identity matrix
I = np.array([[1,0],[0,1]])

epsilon = (e0+e1)*0.5; omega = (e0-e1)*0.5
c = 0.0; omega_z=Xdiag; omega_x = Xnondiag
Hamiltonian = (epsilon+c)*I+(omega_z+omega)*Z+omega_x*X
EigValues, EigVectors = np.linalg.eig(Hamiltonian)
permute = EigValues.argsort()
EigValues = EigValues[permute]
# print only the lowest eigenvalue
print(EigValues[0])
```

Implementing the VQE

For a one-qubit system we can reach every point on the Bloch sphere (as discussed earlier) with a rotation about the x -axis and the y -axis.

We can express this mathematically through the following operations (see whiteboard for the drawing), giving us a new state $|\psi\rangle$

$$|\psi\rangle = R_y(\phi)R_x(\theta)|0\rangle.$$

We can produce multiple ansatzes for the new state in terms of the angles θ and ϕ . With these ansatzes we can in turn calculate the expectation value of the above Hamiltonian, now rewritten in terms of various Pauli matrices (and thereby gates), that is compute

$$\langle \psi | (c + \mathcal{E})\mathbf{I} + (\Omega + \omega_z)\boldsymbol{\sigma}_z + \omega_x\boldsymbol{\sigma}_x | \psi \rangle.$$

We can now set up a series of ansatzes for $|\psi\rangle$ as function of the angles θ and ϕ and find thereafter the variational minimum using for example a gradient descent method.

To do so, we need to remind ourselves about the mathematical expressions for the rotational matrices/operators.

$$R_x(\theta) = \cos \frac{\theta}{2} \mathbf{I} - i \sin \frac{\theta}{2} \boldsymbol{\sigma}_x,$$

and

$$R_y(\phi) = \cos \frac{\phi}{2} \mathbf{I} - i \sin \frac{\phi}{2} \boldsymbol{\sigma}_y.$$

```
# define the rotation matrices
# Define angles theta and phi
theta = 0.5*np.pi; phi = 0.2*np.pi
Rx = np.cos(theta*0.5)*I-1j*np.sin(theta*0.5)*X
Ry = np.cos(phi*0.5)*I-1j*np.sin(phi*0.5)*Y
#define basis states
basis0 = np.array([1,0])
basis1 = np.array([0,1])

NewBasis = Ry @ Rx @ basis0
print(NewBasis)
# Compute the expectation value
#Note hermitian conjugation
Energy = NewBasis.conj().T @ Hamiltonian @ NewBasis
print(Energy)
```

Not an impressive results. We set up now a loop over many angles θ and ϕ and compute the energies

```
# define a number of angles
n = 20
angle = np.arange(0,180,10)
n = np.size(angle)
ExpectationValues = np.zeros((n,n))
for i in range (n):
    theta = np.pi*angle[i]/180.0
    Rx = np.cos(theta*0.5)*I-1j*np.sin(theta*0.5)*X
    for j in range (n):
        phi = np.pi*angle[j]/180.0
        Ry = np.cos(phi*0.5)*I-1j*np.sin(phi*0.5)*Y
        NewBasis = Ry @ Rx @ basis0
        Energy = NewBasis.conj().T @ Hamiltonian @ NewBasis
        Edifference=abs(np.real(EigValues[0])-Energy)
        ExpectationValues[i,j]=Edifference

print(np.min(ExpectationValues))
```

Clearly, this is not the best way of proceeding. Rather, here we could try to find the optimal values for the parameters θ and ϕ through computation of their respective gradients and thereby find the minimum as function of the optimal angles $\hat{\theta}$ and $\hat{\phi}$.

Let us now implement a classical gradient descent algorithm to the computation of the energies. We will follow closely <https://journals.aps.org/prabstract/10.1103/PhysRevA.99.032331> in order to calculate gradients of the Hamiltonian.

Gradient descent and calculations of gradients

In order to optimize the VQE ansatz, we need to compute derivatives with respect to the variational parameters. Here we develop first a simpler approach tailored to the one-qubit case. For this particular case, we have defined an ansatz in terms of the Pauli rotation matrices. These define an arbitrary one-qubit state on the Bloch sphere through the expression

$$|\psi\rangle = |\psi(\theta, \phi)\rangle = R_y(\phi)R_x(\theta)|0\rangle.$$

Each of these rotation matrices can be written in a more general form as

$$R_i(\gamma) = \exp(-i\frac{\gamma}{2}\sigma_i) = \cos(\frac{\gamma}{2})\mathbf{I} - i\sin(\frac{\gamma}{2})\sigma_i,$$

where σ_i is one of the Pauli matrices $\sigma_{x,y,z}$.

It is easy to see that the derivative with respect to γ is

$$\frac{\partial R_i(\gamma)}{\partial \gamma} = -\frac{\gamma}{2}\sigma_i R_i(\gamma).$$

We can now calculate the derivative of the expectation value of the Hamiltonian in terms of the angles θ and ϕ . We have two derivatives

$$\frac{\partial}{\partial \theta} [\langle \psi(\theta, \phi) | \mathbf{H} | \psi(\theta, \phi) \rangle] = \frac{\partial}{\partial \theta} [\langle \mathbf{H}(\theta, \phi) \rangle] = \langle \psi(\theta, \phi) | \mathbf{H}(-\frac{i}{2}\sigma_x | \psi(\theta, \phi) \rangle + \text{h.c.},$$

and

$$\frac{\partial}{\partial \phi} [\langle \psi(\theta, \phi) | \mathbf{H} | \psi(\theta, \phi) \rangle] = \frac{\partial}{\partial \phi} [\langle \mathbf{H}(\theta, \phi) \rangle] = \langle \psi(\theta, \phi) | \mathbf{H}(-\frac{i}{2}\sigma_y | \psi(\theta, \phi) \rangle + \text{h.c.}$$

This means that we have to calculate two additional expectation values in addition to the expectation value of the Hamiltonian itself. If we stay with an ansatz for the single qubit states given by the above rotation operators, we can, following for example [the article by Maria Schuld et al](#), show that the derivative of the expectation value of the Hamiltonian can be written as (we focus only on a given angle ϕ)

$$\frac{\partial}{\partial \phi} [\langle \mathbf{H}(\phi) \rangle] = \frac{1}{2} \left[\langle \mathbf{H}(\phi + \frac{\pi}{2}) \rangle - \langle \mathbf{H}(\phi - \frac{\pi}{2}) \rangle \right].$$

To see this, consider again the definition of the rotation operators. We can write these operators as

$$R_i(\phi) = \exp -\imath(\phi\sigma_i),$$

with ***sigma***_{*i*}, with σ_i being any of the Pauli matrices *X*, *Y* and *Z*. The latter can be generalized to other unitary matrices as well. The derivative with respect to ϕ gives

$$\frac{\partial R_i(\phi)}{\partial \phi} = -\frac{\imath}{2}\sigma_i \exp -\imath(\phi\sigma_i) = -\frac{\imath}{2}\sigma R_i(\phi).$$

Our ansatz for a general one-qubit state on the Bloch sphere contains the product of a rotation around the *x*-axis and the *y*-axis. In the derivation here we focus only on one angle however. Our ansatz is then given by

$$|\psi\rangle = R_i(\phi)|0\rangle,$$

and the expectation value of our Hamiltonian is

$$\langle\psi|\hat{H}|\psi\rangle = \langle 0|R_i(\phi)^\dagger \hat{H} R_i(\phi)|0\rangle.$$

Our derivative with respect to the angle ϕ has a similar structure, that is

$$\frac{\partial}{\partial \phi} [\langle\psi(\theta, \phi)|\mathbf{H}|\psi(\theta, \phi)\rangle] = \langle\psi(\theta, \phi)|\mathbf{H}(-\frac{\imath}{2}\sigma_y|\psi(\theta, \phi)\rangle + \text{h.c.})$$

In order to rewrite the equation of the derivative, the following relation is useful

$$\langle\psi|\hat{A}^\dagger \hat{B} \hat{C}|\psi\rangle = \frac{1}{2} \left[\langle\psi|(\hat{A} + \hat{C})^\dagger \hat{B}(\mathbf{A} + \hat{C})|\psi\rangle - \langle\psi|(\hat{A} - \hat{C})^\dagger \hat{B}(\mathbf{A} - \hat{C})|\psi\rangle \right],$$

where \hat{A} , \hat{B} and \hat{C} are arbitrary hermitian operators. If we identify these operators as $\hat{A} = \mathbf{I}$, with \mathbf{I} being the unit operator, $\hat{B} = \hat{H}$ our Hamiltonian, and $\hat{C} = -\imath\sigma_i/2$, we obtain the following expression for the expectation value of the derivative (excluding the hermitian conjugate)

$$\langle\psi|\mathbf{I}^\dagger \hat{H}(-\frac{\imath}{2}\sigma_i|\psi\rangle = \frac{1}{2} \left[\langle\psi|(\mathbf{I} - \frac{\imath}{2}\sigma_i)^\dagger \hat{H}(\mathbf{I} - \frac{\imath}{2}\sigma_i)|\psi\rangle - \langle\psi|(\mathbf{I} + \frac{\imath}{2}\sigma_i)^\dagger \hat{H}(\mathbf{I} + \frac{\imath}{2}\sigma_i)|\psi\rangle \right].$$

If we then use that the rotation matrices can be rewritten as

$$R_i(\phi) = \exp -(\imath\frac{\phi}{2}\sigma_i) = \cos(\frac{\phi}{2})\mathbf{I} - \imath\sin(\frac{\phi}{2})\sigma_i,$$

we see that if we set the angle to $\phi = \pi/2$, we have

$$R_i(\frac{\pi}{2}) = \cos(\frac{\pi}{4})\mathbf{I} - \imath\sin(\frac{\pi}{4})\sigma_i = \frac{1}{\sqrt{2}} \left(\mathbf{I} - \frac{\imath}{2}\sigma_i \right).$$

This means that we can write

$$\langle\psi|\mathbf{I}^\dagger \hat{H}(-\frac{\imath}{2}\sigma_i|\psi\rangle = \frac{1}{2} \left[\langle\psi|R_i(\frac{\pi}{2})^\dagger \hat{H} R_i(\frac{\pi}{2})|\psi\rangle - \langle\psi|R_i(-\frac{\pi}{2})^\dagger \hat{H} R_i(-\frac{\pi}{2})|\psi\rangle \right] = \frac{1}{2} (\langle\hat{H}(\phi+\frac{\pi}{2})\rangle - \langle\hat{H}(\phi-\frac{\pi}{2})\rangle).$$

Basics of gradient descent and stochastic gradient descent

In order to implement the above equations, we need to remind the reader about basic elements of various optimization approaches. Our main focus here will be various gradient descent approaches, with and without momentum and various approaches for approximating the second derivatives via adaptive methods.

The material here is covered by the lectures from [FYS-STK4155 on gradient optimization](#).

We will repeat some of the basic ingredients from these lectures.

Brief reminder on Newton-Raphson's method

Let us quickly remind ourselves how we derive the above method.

Perhaps the most celebrated of all one-dimensional root-finding routines is Newton's method, also called the Newton-Raphson method. This method requires the evaluation of both the function f and its derivative f' at arbitrary points. If you can only calculate the derivative numerically and/or your function is not of the smooth type, we normally discourage the use of this method.

The equations

The Newton-Raphson formula consists geometrically of extending the tangent line at a current point until it crosses zero, then setting the next guess to the abscissa of that zero-crossing. The mathematics behind this method is rather simple. Employing a Taylor expansion for x sufficiently close to the solution s , we have

$$f(s) = 0 = f(x) + (s - x)f'(x) + \frac{(s - x)^2}{2}f''(x) + \dots$$

For small enough values of the function and for well-behaved functions, the terms beyond linear are unimportant, hence we obtain

$$f(x) + (s - x)f'(x) \approx 0,$$

yielding

$$s \approx x - \frac{f(x)}{f'(x)}.$$

Having in mind an iterative procedure, it is natural to start iterating with

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}.$$

Simple geometric interpretation

The above is Newton-Raphson's method. It has a simple geometric interpretation, namely x_{n+1} is the point where the tangent from $(x_n, f(x_n))$ crosses the x -axis. Close to the solution, Newton-Raphson converges fast to the desired result.

However, if we are far from a root, where the higher-order terms in the series are important, the Newton-Raphson formula can give grossly inaccurate results. For instance, the initial guess for the root might be so far from the true root as to let the search interval include a local maximum or minimum of the function. If an iteration places a trial guess near such a local extremum, so that the first derivative nearly vanishes, then Newton-Raphson may fail totally

Extending to more than one variable

Newton's method can be generalized to systems of several non-linear equations and variables. Consider the case with two equations

$$\begin{aligned} f_1(x_1, x_2) &= 0 \\ f_2(x_1, x_2) &= 0, \end{aligned}$$

which we Taylor expand to obtain

$$\begin{aligned} 0 = f_1(x_1 + h_1, x_2 + h_2) &= f_1(x_1, x_2) + h_1 \partial f_1 / \partial x_1 + h_2 \partial f_1 / \partial x_2 + \dots \\ 0 = f_2(x_1 + h_1, x_2 + h_2) &= f_2(x_1, x_2) + h_1 \partial f_2 / \partial x_1 + h_2 \partial f_2 / \partial x_2 + \dots \end{aligned}$$

Defining the Jacobian matrix \mathbf{J} we have

$$\mathbf{J} = \begin{pmatrix} \partial f_1 / \partial x_1 & \partial f_1 / \partial x_2 \\ \partial f_2 / \partial x_1 & \partial f_2 / \partial x_2 \end{pmatrix},$$

we can rephrase Newton's method as

$$\begin{pmatrix} x_1^{n+1} \\ x_2^{n+1} \end{pmatrix} = \begin{pmatrix} x_1^n \\ x_2^n \end{pmatrix} + \begin{pmatrix} h_1^n \\ h_2^n \end{pmatrix},$$

where we have defined

$$\begin{pmatrix} h_1^n \\ h_2^n \end{pmatrix} = -\mathbf{J}^{-1} \begin{pmatrix} f_1(x_1^n, x_2^n) \\ f_2(x_1^n, x_2^n) \end{pmatrix}.$$

We need thus to compute the inverse of the Jacobian matrix and it is to understand that difficulties may arise in case \mathbf{J} is nearly singular.

It is rather straightforward to extend the above scheme to systems of more than two non-linear equations. In our case, the Jacobian matrix is given by the Hessian that represents the second derivative of cost function.

Steepest descent

The basic idea of gradient descent is that a function $F(\mathbf{x})$, $\mathbf{x} \equiv (x_1, \dots, x_n)$, decreases fastest if one goes from \mathbf{x} in the direction of the negative gradient $-\nabla F(\mathbf{x})$.

It can be shown that if

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \gamma_k \nabla F(\mathbf{x}_k),$$

with $\gamma_k > 0$.

For γ_k small enough, then $F(\mathbf{x}_{k+1}) \leq F(\mathbf{x}_k)$. This means that for a sufficiently small γ_k we are always moving towards smaller function values, i.e a minimum.

More on Steepest descent

The previous observation is the basis of the method of steepest descent, which is also referred to as just gradient descent (GD). One starts with an initial guess \mathbf{x}_0 for a minimum of F and computes new approximations according to

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \gamma_k \nabla F(\mathbf{x}_k), \quad k \geq 0.$$

The parameter γ_k is often referred to as the step length or the learning rate within the context of Machine Learning.

The ideal

Ideally the sequence $\{\mathbf{x}_k\}_{k=0}$ converges to a global minimum of the function F . In general we do not know if we are in a global or local minimum. In the special case when F is a convex function, all local minima are also global minima, so in this case gradient descent can converge to the global solution. The advantage of this scheme is that it is conceptually simple and straightforward to implement. However the method in this form has some severe limitations:

In machine learning we are often faced with non-convex high dimensional cost functions with many local minima. Since GD is deterministic we will get stuck in a local minimum, if the method converges, unless we have a very good initial guess. This also implies that the scheme is sensitive to the chosen initial condition.

Note that the gradient is a function of $\mathbf{x} = (x_1, \dots, x_n)$ which makes it expensive to compute numerically.

The sensitiveness of the gradient descent

The gradient descent method is sensitive to the choice of learning rate γ_k . This is due to the fact that we are only guaranteed that $F(\mathbf{x}_{k+1}) \leq F(\mathbf{x}_k)$ for sufficiently small γ_k . The problem is to determine an optimal learning rate. If the learning rate is chosen too small the method will take a long time to converge and if it is too large we can experience erratic behavior.

Many of these shortcomings can be alleviated by introducing randomness. One such method is that of Stochastic Gradient Descent (SGD), see below.

Convex functions

Ideally we want our cost/loss function to be convex(concave).

First we give the definition of a convex set: A set C in \mathbb{R}^n is said to be convex if, for all x and y in C and all $t \in (0, 1)$, the point $(1 - t)x + ty$ also belongs to

C. Geometrically this means that every point on the line segment connecting x and y is in C as discussed below.

The convex subsets of \mathbb{R} are the intervals of \mathbb{R} . Examples of convex sets of \mathbb{R}^2 are the regular polygons (triangles, rectangles, pentagons, etc...).

Convex function

Convex function: Let $X \subset \mathbb{R}^n$ be a convex set. Assume that the function $f : X \rightarrow \mathbb{R}$ is continuous, then f is said to be convex if

$$f(tx_1 + (1-t)x_2) \leq tf(x_1) + (1-t)f(x_2)$$

for all $x_1, x_2 \in X$ and for all $t \in [0, 1]$. If \leq is replaced with a strict inequality in the definition, we demand $x_1 \neq x_2$ and $t \in (0, 1)$ then f is said to be strictly convex. For a single variable function, convexity means that if you draw a straight line connecting $f(x_1)$ and $f(x_2)$, the value of the function on the interval $[x_1, x_2]$ is always below the line as illustrated below.

Conditions on convex functions

In the following we state first and second-order conditions which ensures convexity of a function f . We write D_f to denote the domain of f , i.e the subset of \mathbb{R}^n where f is defined. For more details and proofs we refer to: [S. Boyd and L. Vandenberghe. Convex Optimization. Cambridge University Press.](#)

First order condition. Suppose f is differentiable (i.e $\nabla f(x)$ is well defined for all x in the domain of f). Then f is convex if and only if D_f is a convex set and

$$f(y) \geq f(x) + \nabla f(x)^T(y - x)$$

holds for all $x, y \in D_f$. This condition means that for a convex function the first order Taylor expansion (right hand side above) at any point is a global under estimator of the function. To convince yourself you can make a drawing of $f(x) = x^2 + 1$ and draw the tangent line to $f(x)$ and note that it is always below the graph.

Second order condition. Assume that f is twice differentiable, i.e the Hessian matrix exists at each point in D_f . Then f is convex if and only if D_f is a convex set and its Hessian is positive semi-definite for all $x \in D_f$. For a single-variable function this reduces to $f''(x) \geq 0$. Geometrically this means that f has nonnegative curvature everywhere.

This condition is particularly useful since it gives us an procedure for determining if the function under consideration is convex, apart from using the definition.

More on convex functions

The next result is of great importance to us and the reason why we are going on about convex functions. In machine learning we frequently have to minimize a loss/cost function in order to find the best parameters for the model we are considering.

Ideally we want the global minimum (for high-dimensional models it is hard to know if we have local or global minimum). However, if the cost/loss function is convex the following result provides invaluable information:

Any minimum is global for convex functions. Consider the problem of finding $x \in \mathbb{R}^n$ such that $f(x)$ is minimal, where f is convex and differentiable. Then, any point x^* that satisfies $\nabla f(x^*) = 0$ is a global minimum.

This result means that if we know that the cost/loss function is convex and we are able to find a minimum, we are guaranteed that it is a global minimum.

Some simple problems

1. Show that $f(x) = x^2$ is convex for $x \in \mathbb{R}$ using the definition of convexity.
Hint: If you re-write the definition, f is convex if the following holds for all $x, y \in D_f$ and any $\lambda \in [0, 1]$ $\lambda f(x) + (1 - \lambda)f(y) - f(\lambda x + (1 - \lambda)y) \geq 0$.
2. Using the second order condition show that the following functions are convex on the specified domain.
 - $f(x) = e^x$ is convex for $x \in \mathbb{R}$.
 - $g(x) = -\ln(x)$ is convex for $x \in (0, \infty)$.
3. Let $f(x) = x^2$ and $g(x) = e^x$. Show that $f(g(x))$ and $g(f(x))$ is convex for $x \in \mathbb{R}$. Also show that if $f(x)$ is any convex function then $h(x) = e^{f(x)}$ is convex.
4. A norm is any function that satisfy the following properties
 - $f(\alpha x) = |\alpha|f(x)$ for all $\alpha \in \mathbb{R}$.
 - $f(x + y) \leq f(x) + f(y)$
 - $f(x) \leq 0$ for all $x \in \mathbb{R}^n$ with equality if and only if $x = 0$

Using the definition of convexity, try to show that a function satisfying the properties above is convex (the third condition is not needed to show this).

Gradient descent example

Let $\mathbf{y} = (y_1, \dots, y_n)^T$, $\mathbf{y} = (\mathbf{y}_1, \dots, \mathbf{y}_n)^T$ and $\beta = (\beta_0, \beta_1)^T$

It is convenient to write $\mathbf{y} = X\beta$ where $X \in \mathbb{R}^{100 \times 2}$ is the design matrix given by (we keep the intercept here)

$$X \equiv \begin{bmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_{100} \end{bmatrix}.$$

The cost/loss/risk function is given by (

$$C(\beta) = \frac{1}{n} \|X\beta - \mathbf{y}\|_2^2 = \frac{1}{n} \sum_{i=1}^{100} [(\beta_0 + \beta_1 x_i)^2 - 2y_i(\beta_0 + \beta_1 x_i) + y_i^2]$$

and we want to find β such that $C(\beta)$ is minimized.

The derivative of the cost/loss function

Computing $\partial C(\beta)/\partial \beta_0$ and $\partial C(\beta)/\partial \beta_1$ we can show that the gradient can be written as

$$\nabla_{\beta} C(\beta) = \frac{2}{n} \left[\begin{array}{c} \sum_{i=1}^{100} (\beta_0 + \beta_1 x_i - y_i) \\ \sum_{i=1}^{100} (x_i (\beta_0 + \beta_1 x_i) - y_i x_i) \end{array} \right] = \frac{2}{n} X^T (X\beta - \mathbf{y}),$$

where X is the design matrix defined above.

The Hessian matrix

The Hessian matrix of $C(\beta)$ is given by

$$\mathbf{H} \equiv \begin{bmatrix} \frac{\partial^2 C(\beta)}{\partial \beta_0^2} & \frac{\partial^2 C(\beta)}{\partial \beta_0 \partial \beta_1} \\ \frac{\partial^2 C(\beta)}{\partial \beta_0 \partial \beta_1} & \frac{\partial^2 C(\beta)}{\partial \beta_1^2} \end{bmatrix} = \frac{2}{n} X^T X.$$

This result implies that $C(\beta)$ is a convex function since the matrix $X^T X$ always is positive semi-definite.

Simple program

We can now write a program that minimizes $C(\beta)$ using the gradient descent method with a constant learning rate γ according to

$$\beta_{k+1} = \beta_k - \gamma \nabla_{\beta} C(\beta_k), \quad k = 0, 1, \dots$$

We can use the expression we computed for the gradient and let use a β_0 be chosen randomly and let $\gamma = 0.001$. Stop iterating when $\|\nabla_{\beta} C(\beta_k)\| \leq \epsilon = 10^{-8}$.

Note that the code below does not include the latter stop criterion.

And finally we can compare our solution for β with the analytic result given by $\beta = (X^T X)^{-1} X^T \mathbf{y}$.

Gradient Descent Example

Here our simple example

```
# Importing various packages
from random import random, seed
import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from matplotlib import cm
from matplotlib.ticker import LinearLocator, FormatStrFormatter
import sys

# the number of datapoints
n = 100
x = 2*np.random.rand(n,1)
y = 4+3*x+np.random.randn(n,1)

X = np.c_[np.ones((n,1)), x]
# Hessian matrix
H = (2.0/n)* X.T @ X
# Get the eigenvalues
EigValues, EigVectors = np.linalg.eig(H)
print(f"Eigenvalues of Hessian Matrix:{EigValues}")

beta_linreg = np.linalg.inv(X.T @ X) @ X.T @ y
print(beta_linreg)
beta = np.random.randn(2,1)

eta = 1.0/np.max(EigValues)
Niterations = 1000

for iter in range(Niterations):
    gradient = (2.0/n)*X.T @ (X @ beta-y)
    beta -= eta*gradient

print(beta)
xnew = np.array([[0],[2]])
xbnew = np.c_[np.ones((2,1)), xnew]
ypredict = xbnew.dot(beta)
ypredict2 = xbnew.dot(beta_linreg)
plt.plot(xnew, ypredict, "r-")
plt.plot(xnew, ypredict2, "b-")
plt.plot(x, y, 'ro')
plt.axis([0,2.0,0, 15.0])
plt.xlabel(r'$x$')
plt.ylabel(r'$y$')
plt.title(r'Gradient descent example')
plt.show()
```

Using gradient descent methods, limitations

- **Gradient descent (GD) finds local minima of our function.** Since the GD algorithm is deterministic, if it converges, it will converge to a local minimum of our cost/loss/risk function. Because in ML we are often

dealing with extremely rugged landscapes with many local minima, this can lead to poor performance.

- **GD is sensitive to initial conditions.** One consequence of the local nature of GD is that initial conditions matter. Depending on where one starts, one will end up at a different local minima. Therefore, it is very important to think about how one initializes the training process. This is true for GD as well as more complicated variants of GD.
- **Gradients are computationally expensive to calculate for large datasets.** In many cases in statistics and ML, the cost/loss/risk function is a sum of terms, with one term for each data point. For example, in linear regression, $E \propto \sum_{i=1}^n (y_i - \mathbf{w}^T \cdot \mathbf{x}_i)^2$; for logistic regression, the square error is replaced by the cross entropy. To calculate the gradient we have to sum over *all* n data points. Doing this at every GD step becomes extremely computationally expensive. An ingenious solution to this, is to calculate the gradients using small subsets of the data called “mini batches”. This has the added benefit of introducing stochasticity into our algorithm.
- **GD is very sensitive to choices of learning rates.** GD is extremely sensitive to the choice of learning rates. If the learning rate is very small, the training process take an extremely long time. For larger learning rates, GD can diverge and give poor results. Furthermore, depending on what the local landscape looks like, we have to modify the learning rates to ensure convergence. Ideally, we would *adaptively* choose the learning rates to match the landscape.
- **GD treats all directions in parameter space uniformly.** Another major drawback of GD is that unlike Newton’s method, the learning rate for GD is the same in all directions in parameter space. For this reason, the maximum learning rate is set by the behavior of the steepest direction and this can significantly slow down training. Ideally, we would like to take large steps in flat directions and small steps in steep directions. Since we are exploring rugged landscapes where curvatures change, this requires us to keep track of not only the gradient but second derivatives. The ideal scenario would be to calculate the Hessian but this proves to be too computationally expensive.
- GD can take exponential time to escape saddle points, even with random initialization. As we mentioned, GD is extremely sensitive to initial condition since it determines the particular local minimum GD would eventually reach. However, even with a good initialization scheme, through the introduction of randomness, GD can still take exponential time to escape saddle points.

Improving gradient descent with momentum

We discuss here some simple examples where we introduce what is called 'memory' about previous steps, or what is normally called momentum gradient descent. The mathematics is explained below in connection with Stochastic gradient descent.

```
from numpy import asarray
from numpy import arange
from numpy.random import rand
from numpy.random import seed
from matplotlib import pyplot

# objective function
def objective(x):
    return x**2.0

# derivative of objective function
def derivative(x):
    return x * 2.0

# gradient descent algorithm
def gradient_descent(objective, derivative, bounds, n_iter, step_size):
    # track all solutions
    solutions, scores = list(), list()
    # generate an initial point
    solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
    # run the gradient descent
    for i in range(n_iter):
        # calculate gradient
        gradient = derivative(solution)
        # take a step
        solution = solution - step_size * gradient
        # evaluate candidate point
        solution_eval = objective(solution)
        # store solution
        solutions.append(solution)
        scores.append(solution_eval)
        # report progress
        print('>%d f(%s) = %.5f' % (i, solution, solution_eval))
    return [solutions, scores]

# seed the pseudo random number generator
seed(4)
# define range for input
bounds = asarray([[-1.0, 1.0]])
# define the total iterations
n_iter = 30
# define the step size
step_size = 0.1
# perform the gradient descent search
solutions, scores = gradient_descent(objective, derivative, bounds, n_iter, step_size)
# sample input range uniformly at 0.1 increments
inputs = arange(bounds[0,0], bounds[0,1]+0.1, 0.1)
# compute targets
results = objective(inputs)
# create a line plot of input vs result
pyplot.plot(inputs, results)
# plot the solutions found
pyplot.plot(solutions, scores, '.-', color='red')
```

```

# show the plot
pyplot.show()

```

Same code but now with momentum gradient descent

```

from numpy import asarray
from numpy import arange
from numpy.random import rand
from numpy.random import seed
from matplotlib import pyplot

# objective function
def objective(x):
    return x**2.0

# derivative of objective function
def derivative(x):
    return x * 2.0

# gradient descent algorithm
def gradient_descent(objective, derivative, bounds, n_iter, step_size, momentum):
    # track all solutions
    solutions, scores = list(), list()
    # generate an initial point
    solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
    # keep track of the change
    change = 0.0
    # run the gradient descent
    for i in range(n_iter):
        # calculate gradient
        gradient = derivative(solution)
        # calculate update
        new_change = step_size * gradient + momentum * change
        # take a step
        solution = solution - new_change
        # save the change
        change = new_change
        # evaluate candidate point
        solution_eval = objective(solution)
        # store solution
        solutions.append(solution)
        scores.append(solution_eval)
        # report progress
        print('>%d f(%s) = %.5f' % (i, solution, solution_eval))
    return [solutions, scores]

# seed the pseudo random number generator
seed(4)
# define range for input
bounds = asarray([[-1.0, 1.0]])
# define the total iterations
n_iter = 30
# define the step size
step_size = 0.1
# define momentum
momentum = 0.3
# perform the gradient descent search with momentum
solutions, scores = gradient_descent(objective, derivative, bounds, n_iter, step_size, momentum)

```

```

# sample input range uniformly at 0.1 increments
inputs = arange(bounds[0,0], bounds[0,1]+0.1, 0.1)
# compute targets
results = objective(inputs)
# create a line plot of input vs result
pyplot.plot(inputs, results)
# plot the solutions found
pyplot.plot(solutions, scores, '.-', color='red')
# show the plot
pyplot.show()

```

Momentum based GD

The stochastic gradient descent (SGD) is almost always used with a *momentum* or inertia term that serves as a memory of the direction we are moving in parameter space. This is typically implemented as follows

$$\begin{aligned}\mathbf{v}_t &= \gamma \mathbf{v}_{t-1} + \eta_t \nabla_{\theta} E(\boldsymbol{\theta}_t) \\ \boldsymbol{\theta}_{t+1} &= \boldsymbol{\theta}_t - \mathbf{v}_t,\end{aligned}\tag{63}$$

where we have introduced a momentum parameter γ , with $0 \leq \gamma \leq 1$, and for brevity we dropped the explicit notation to indicate the gradient is to be taken over a different mini-batch at each step. We call this algorithm gradient descent with momentum (GDM). From these equations, it is clear that \mathbf{v}_t is a running average of recently encountered gradients and $(1 - \gamma)^{-1}$ sets the characteristic time scale for the memory used in the averaging procedure. Consistent with this, when $\gamma = 0$, this just reduces down to ordinary SGD as discussed earlier. An equivalent way of writing the updates is

$$\Delta \boldsymbol{\theta}_{t+1} = \gamma \Delta \boldsymbol{\theta}_t - \eta_t \nabla_{\theta} E(\boldsymbol{\theta}_t),$$

where we have defined $\Delta \boldsymbol{\theta}_t = \boldsymbol{\theta}_t - \boldsymbol{\theta}_{t-1}$.

More on momentum based approaches

Let us try to get more intuition from these equations. It is helpful to consider a simple physical analogy with a particle of mass m moving in a viscous medium with drag coefficient μ and potential $E(\mathbf{w})$. If we denote the particle's position by \mathbf{w} , then its motion is described by

$$m \frac{d^2 \mathbf{w}}{dt^2} + \mu \frac{d\mathbf{w}}{dt} = -\nabla_{\mathbf{w}} E(\mathbf{w}).$$

We can discretize this equation in the usual way to get

$$m \frac{\mathbf{w}_{t+\Delta t} - 2\mathbf{w}_t + \mathbf{w}_{t-\Delta t}}{(\Delta t)^2} + \mu \frac{\mathbf{w}_{t+\Delta t} - \mathbf{w}_t}{\Delta t} = -\nabla_{\mathbf{w}} E(\mathbf{w}).$$

Rearranging this equation, we can rewrite this as

$$\Delta \mathbf{w}_{t+\Delta t} = -\frac{(\Delta t)^2}{m + \mu \Delta t} \nabla_w E(\mathbf{w}) + \frac{m}{m + \mu \Delta t} \Delta \mathbf{w}_t.$$

Momentum parameter

Notice that this equation is identical to previous one if we identify the position of the particle, \mathbf{w} , with the parameters $\boldsymbol{\theta}$. This allows us to identify the momentum parameter and learning rate with the mass of the particle and the viscous drag as:

$$\gamma = \frac{m}{m + \mu \Delta t}, \quad \eta = \frac{(\Delta t)^2}{m + \mu \Delta t}.$$

Thus, as the name suggests, the momentum parameter is proportional to the mass of the particle and effectively provides inertia. Furthermore, in the large viscosity/small learning rate limit, our memory time scales as $(1 - \gamma)^{-1} \approx m/(\mu \Delta t)$.

Why is momentum useful? SGD momentum helps the gradient descent algorithm gain speed in directions with persistent but small gradients even in the presence of stochasticity, while suppressing oscillations in high-curvature directions. This becomes especially important in situations where the landscape is shallow and flat in some directions and narrow and steep in others. It has been argued that first-order methods (with appropriate initial conditions) can perform comparable to more expensive second order methods, especially in the context of complex deep learning models.

These beneficial properties of momentum can sometimes become even more pronounced by using a slight modification of the classical momentum algorithm called Nesterov Accelerated Gradient (NAG).

In the NAG algorithm, rather than calculating the gradient at the current parameters, $\nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}_t)$, one calculates the gradient at the expected value of the parameters given our current momentum, $\nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}_t + \gamma \mathbf{v}_{t-1})$. This yields the NAG update rule

$$\begin{aligned} \mathbf{v}_t &= \gamma \mathbf{v}_{t-1} + \eta_t \nabla_{\boldsymbol{\theta}} E(\boldsymbol{\theta}_t + \gamma \mathbf{v}_{t-1}) \\ \boldsymbol{\theta}_{t+1} &= \boldsymbol{\theta}_t - \mathbf{v}_t. \end{aligned} \tag{64}$$

One of the major advantages of NAG is that it allows for the use of a larger learning rate than GDM for the same choice of γ .

Second moment of the gradient

In stochastic gradient descent, with and without momentum, we still have to specify a schedule for tuning the learning rates η_t as a function of time. As discussed in the context of Newton's method, this presents a number of dilemmas. The learning rate is limited by the steepest direction which can change depending

on the current position in the landscape. To circumvent this problem, ideally our algorithm would keep track of curvature and take large steps in shallow, flat directions and small steps in steep, narrow directions. Second-order methods accomplish this by calculating or approximating the Hessian and normalizing the learning rate by the curvature. However, this is very computationally expensive for extremely large models. Ideally, we would like to be able to adaptively change the step size to match the landscape without paying the steep computational price of calculating or approximating Hessians.

Recently, a number of methods have been introduced that accomplish this by tracking not only the gradient, but also the second moment of the gradient. These methods include AdaGrad, AdaDelta, Root Mean Squared Propagation (RMS-Prop), and [ADAM](#).

RMS prop

In RMS prop, in addition to keeping a running average of the first moment of the gradient, we also keep track of the second moment denoted by $\mathbf{s}_t = \mathbb{E}[\mathbf{g}_t^2]$. The update rule for RMS prop is given by

$$\begin{aligned} \mathbf{g}_t &= \nabla_{\theta} E(\boldsymbol{\theta}) \\ \mathbf{s}_t &= \beta \mathbf{s}_{t-1} + (1 - \beta) \mathbf{g}_t^2 \\ \boldsymbol{\theta}_{t+1} &= \boldsymbol{\theta}_t - \eta_t \frac{\mathbf{g}_t}{\sqrt{\mathbf{s}_t + \epsilon}}, \end{aligned} \tag{65}$$

where β controls the averaging time of the second moment and is typically taken to be about $\beta = 0.9$, η_t is a learning rate typically chosen to be 10^{-3} , and $\epsilon \sim 10^{-8}$ is a small regularization constant to prevent divergences. Multiplication and division by vectors is understood as an element-wise operation. It is clear from this formula that the learning rate is reduced in directions where the norm of the gradient is consistently large. This greatly speeds up the convergence by allowing us to use a larger learning rate for flat directions.

ADAM optimizer

A related algorithm is the ADAM optimizer. In [ADAM](#), we keep a running average of both the first and second moment of the gradient and use this information to adaptively change the learning rate for different parameters. The method is efficient when working with large problems involving lots of data and/or parameters. It is a combination of the gradient descent with momentum algorithm and the RMSprop algorithm discussed above.

In addition to keeping a running average of the first and second moments of the gradient (i.e. $\mathbf{m}_t = \mathbb{E}[\mathbf{g}_t]$ and $\mathbf{s}_t = \mathbb{E}[\mathbf{g}_t^2]$, respectively), ADAM performs an additional bias correction to account for the fact that we are estimating the first two moments of the gradient using a running average (denoted by the hats in the

update rule below). The update rule for ADAM is given by (where multiplication and division are once again understood to be element-wise operations below)

$$\begin{aligned}
\mathbf{g}_t &= \nabla_{\theta} E(\boldsymbol{\theta}) \\
\mathbf{m}_t &= \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t \\
\mathbf{s}_t &= \beta_2 \mathbf{s}_{t-1} + (1 - \beta_2) \mathbf{g}_t^2 \\
\hat{\mathbf{m}}_t &= \frac{\mathbf{m}_t}{1 - \beta_1^t} \\
\hat{\mathbf{s}}_t &= \frac{\mathbf{s}_t}{1 - \beta_2^t} \\
\boldsymbol{\theta}_{t+1} &= \boldsymbol{\theta}_t - \eta_t \frac{\hat{\mathbf{m}}_t}{\sqrt{\hat{\mathbf{s}}_t} + \epsilon},
\end{aligned} \tag{66}$$

where β_1 and β_2 set the memory lifetime of the first and second moment and are typically taken to be 0.9 and 0.99 respectively, and η and ϵ are identical to RMSprop.

Like in RMSprop, the effective step size of a parameter depends on the magnitude of its gradient squared. To understand this better, let us rewrite this expression in terms of the variance $\sigma_t^2 = \mathbf{s}_t - (\mathbf{m}_t)^2$. Consider a single parameter θ_t . The update rule for this parameter is given by

$$\Delta\theta_{t+1} = -\eta_t \frac{m_t}{\sqrt{\sigma_t^2 + m_t^2} + \epsilon}.$$

Algorithms and codes for Adagrad, RMSprop and Adam

The algorithms we have implemented are well described in the text by [Goodfellow, Bengio and Courville, chapter 8](#).

The codes which implement these algorithms are discussed after our presentation of automatic differentiation.

Practical tips

- **Randomize the data when making mini-batches.** It is always important to randomly shuffle the data when forming mini-batches. Otherwise, the gradient descent method can fit spurious correlations resulting from the order in which data is presented.
- **Transform your inputs.** Learning becomes difficult when our landscape has a mixture of steep and flat directions. One simple trick for minimizing these situations is to standardize the data by subtracting the mean and normalizing the variance of input variables. Whenever possible, also decorrelate the inputs. To understand why this is helpful, consider the case of linear regression. It is easy to show that for the squared error cost

function, the Hessian of the cost function is just the correlation matrix between the inputs. Thus, by standardizing the inputs, we are ensuring that the landscape looks homogeneous in all directions in parameter space. Since most deep networks can be viewed as linear transformations followed by a non-linearity at each layer, we expect this intuition to hold beyond the linear case.

- **Monitor the out-of-sample performance.** Always monitor the performance of your model on a validation set (a small portion of the training data that is held out of the training process to serve as a proxy for the test set. If the validation error starts increasing, then the model is beginning to overfit. Terminate the learning process. This *early stopping* significantly improves performance in many settings.
- **Adaptive optimization methods don't always have good generalization.** Recent studies have shown that adaptive methods such as ADAM, RMSProp, and AdaGrad tend to have poor generalization compared to SGD or SGD with momentum, particularly in the high-dimensional limit (i.e. the number of parameters exceeds the number of data points). Although it is not clear at this stage why these methods perform so well in training deep neural networks, simpler procedures like properly-tuned SGD may work as well or better in these applications.

Using Autograd with OLS

We conclude the part on optimization by showing how we can make codes for linear regression and logistic regression using **autograd**. The first example shows results with ordinary least squares.

```
# Using Autograd to calculate gradients for OLS
from random import random, seed
import numpy as np
import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

def CostOLS(beta):
    return (1.0/n)*np.sum((y-X @ beta)**2)

n = 100
x = 2*np.random.rand(n,1)
y = 4+3*x+np.random.randn(n,1)

X = np.c_[np.ones((n,1)), x]
XT_X = X.T @ X
theta_linreg = np.linalg.pinv(XT_X) @ (X.T @ y)
print("Own inversion")
print(theta_linreg)
# Hessian matrix
H = (2.0/n)* XT_X
EigValues, EigVectors = np.linalg.eig(H)
print(f"Eigenvalues of Hessian Matrix:{EigValues}")
```

```

theta = np.random.randn(2,1)
eta = 1.0/np.max(EigValues)
Niterations = 1000
# define the gradient
training_gradient = grad(CostOLS)

for iter in range(Niterations):
    gradients = training_gradient(theta)
    theta -= eta*gradients
print("theta from own gd")
print(theta)

xnew = np.array([[0],[2]])
Xnew = np.c_[np.ones((2,1)), xnew]
ypredict = Xnew.dot(theta)
ypredict2 = Xnew.dot(theta_linreg)

plt.plot(xnew, ypredict, "r-")
plt.plot(xnew, ypredict2, "b-")
plt.plot(x, y, 'ro')
plt.axis([0,2.0,0, 15.0])
plt.xlabel(r'$x$')
plt.ylabel(r'$y$')
plt.title(r'Random numbers ')
plt.show()

```

Same code but now with momentum gradient descent

```

# Using Autograd to calculate gradients for OLS
from random import random, seed
import numpy as np
import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

def CostOLS(beta):
    return (1.0/n)*np.sum((y-X @ beta)**2)

n = 100
x = 2*np.random.rand(n,1)
y = 4+3*x#+np.random.randn(n,1)

X = np.c_[np.ones((n,1)), x]
XT_X = X.T @ X
theta_linreg = np.linalg.pinv(XT_X) @ (X.T @ y)
print("Own inversion")
print(theta_linreg)
# Hessian matrix
H = (2.0/n)* XT_X
EigValues, EigVectors = np.linalg.eig(H)
print(f"Eigenvalues of Hessian Matrix:{EigValues}")

theta = np.random.randn(2,1)
eta = 1.0/np.max(EigValues)
Niterations = 30

# define the gradient

```

```

training_gradient = grad(CostOLS)

for iter in range(Niterations):
    gradients = training_gradient(theta)
    theta -= eta*gradients
    print(iter,gradients[0],gradients[1])
print("theta from own gd")
print(theta)

# Now improve with momentum gradient descent
change = 0.0
delta_momentum = 0.3
for iter in range(Niterations):
    # calculate gradient
    gradients = training_gradient(theta)
    # calculate update
    new_change = eta*gradients+delta_momentum*change
    # take a step
    theta -= new_change
    # save the change
    change = new_change
    print(iter,gradients[0],gradients[1])
print("theta from own gd with momentum")
print(theta)

```

But none of these can compete with Newton's method

```

# Using Newton's method
from random import random, seed
import numpy as np
import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

def CostOLS(beta):
    return (1.0/n)*np.sum((y-X @ beta)**2)

n = 100
x = 2*np.random.rand(n,1)
y = 4+3*x+np.random.randn(n,1)

X = np.c_[np.ones((n,1)), x]
XT_X = X.T @ X
beta_linreg = np.linalg.pinv(XT_X) @ (X.T @ y)
print("Own inversion")
print(beta_linreg)
# Hessian matrix
H = (2.0/n)* XT_X
# Note that here the Hessian does not depend on the parameters beta
invH = np.linalg.pinv(H)
EigValues, EigVectors = np.linalg.eig(H)
print(f"Eigenvalues of Hessian Matrix:{EigValues}")

beta = np.random.randn(2,1)
Niterations = 5

# define the gradient
training_gradient = grad(CostOLS)

```

```

for iter in range(Niterations):
    gradients = training_gradient(beta)
    beta -= invH @ gradients
    print(iter,gradients[0],gradients[1])
print("beta from own Newton code")
print(beta)

```

Including Stochastic Gradient Descent with Autograd

In this code we include the stochastic gradient descent approach discussed above. Note here that we specify which argument we are taking the derivative with respect to when using **autograd**.

```

# Using Autograd to calculate gradients using SGD
# OLS example
from random import random, seed
import numpy as np
import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

# Note change from previous example
def CostOLS(y,X,theta):
    return np.sum((y-X @ theta)**2)

n = 100
x = 2*np.random.rand(n,1)
y = 4+3*x+np.random.randn(n,1)

X = np.c_[np.ones((n,1)), x]
XT_X = X.T @ X
theta_linreg = np.linalg.pinv(XT_X) @ (X.T @ y)
print("Own inversion")
print(theta_linreg)
# Hessian matrix
H = (2.0/n)* XT_X
EigValues, EigVectors = np.linalg.eig(H)
print(f"Eigenvalues of Hessian Matrix:{EigValues}")

theta = np.random.randn(2,1)
eta = 1.0/np.max(EigValues)
Niterations = 1000

# Note that we request the derivative wrt third argument (theta, 2 here)
training_gradient = grad(CostOLS,2)

for iter in range(Niterations):
    gradients = (1.0/n)*training_gradient(y, X, theta)
    theta -= eta*gradients
    print("theta from own gd")
    print(theta)

xnew = np.array([[0],[2]])
Xnew = np.c_[np.ones((2,1)), xnew]
ypredict = Xnew.dot(theta)
ypredict2 = Xnew.dot(theta_linreg)

```

```

plt.plot(xnew, ypredict, "r-")
plt.plot(xnew, ypredict2, "b-")
plt.plot(x, y, 'ro')
plt.axis([0,2.0,0, 15.0])
plt.xlabel(r'$x$')
plt.ylabel(r'$y$')
plt.title(r'Random numbers ')
plt.show()

n_epochs = 50
M = 5 #size of each minibatch
m = int(n/M) #number of minibatches
t0, t1 = 5, 50
def learning_schedule(t):
    return t0/(t+t1)

theta = np.random.randn(2,1)

for epoch in range(n_epochs):
    # Can you figure out a better way of setting up the contributions to each batch?
    for i in range(m):
        random_index = M*np.random.randint(m)
        xi = X[random_index:random_index+M]
        yi = y[random_index:random_index+M]
        gradients = (1.0/M)*training_gradient(yi, xi, theta)
        eta = learning_schedule(epoch*m+i)
        theta = theta - eta*gradients
print("theta from own sgd")
print(theta)

```

Same code but now with momentum gradient descent

```

# Using Autograd to calculate gradients using SGD
# OLS example
from random import random, seed
import numpy as np
import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

# Note change from previous example
def CostOLS(y,X,theta):
    return np.sum((y-X @ theta)**2)

n = 100
x = 2*np.random.rand(n,1)
y = 4+3*x+np.random.randn(n,1)

X = np.c_[np.ones((n,1)), x]
XT_X = X.T @ X
theta_linreg = np.linalg.pinv(XT_X) @ (X.T @ y)
print("Own inversion")
print(theta_linreg)
# Hessian matrix
H = (2.0/n)* XT_X
EigValues, EigVectors = np.linalg.eig(H)
print(f"Eigenvalues of Hessian Matrix:{EigValues}")

```

```

theta = np.random.randn(2,1)
eta = 1.0/np.max(EigValues)
Niterations = 100

# Note that we request the derivative wrt third argument (theta, 2 here)
training_gradient = grad(CostOLS,2)

for iter in range(Niterations):
    gradients = (1.0/n)*training_gradient(y, X, theta)
    theta -= eta*gradients
    print("theta from own gd")
    print(theta)

n_epochs = 50
M = 5 #size of each minibatch
m = int(n/M) #number of minibatches
t0, t1 = 5, 50
def learning_schedule(t):
    return t0/(t+t1)

theta = np.random.randn(2,1)

change = 0.0
delta_momentum = 0.3

for epoch in range(n_epochs):
    for i in range(m):
        random_index = M*np.random.randint(m)
        xi = X[random_index:random_index+M]
        yi = y[random_index:random_index+M]
        gradients = (1.0/M)*training_gradient(yi, xi, theta)
        eta = learning_schedule(epoch*m+i)
        # calculate update
        new_change = eta*gradients+delta_momentum*change
        # take a step
        theta -= new_change
        # save the change
        change = new_change
    print("theta from own sdg with momentum")
    print(theta)

```

Similar (second order function now) problem but now with AdaGrad

```

# Using Autograd to calculate gradients using AdaGrad and Stochastic Gradient descent
# OLS example
from random import random, seed
import numpy as np
import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

# Note change from previous example
def CostOLS(y,X,theta):
    return np.sum((y-X @ theta)**2)

```

```

n = 1000
x = np.random.rand(n,1)
y = 2.0+3*x +4*x*x

X = np.c_[np.ones((n,1)), x, x*x]
XT_X = X.T @ X
theta_linreg = np.linalg.pinv(XT_X) @ (X.T @ y)
print("Own inversion")
print(theta_linreg)

# Note that we request the derivative wrt third argument (theta, 2 here)
training_gradient = grad(CostOLS,2)
# Define parameters for Stochastic Gradient Descent
n_epochs = 50
M = 5 #size of each minibatch
m = int(n/M) #number of minibatches
# Guess for unknown parameters theta
theta = np.random.randn(3,1)

# Value for learning rate
eta = 0.01
# Including AdaGrad parameter to avoid possible division by zero
delta = 1e-8
for epoch in range(n_epochs):
    Giter = 0.0
    for i in range(m):
        random_index = M*np.random.randint(m)
        xi = X[random_index:random_index+M]
        yi = y[random_index:random_index+M]
        gradients = (1.0/M)*training_gradient(yi, xi, theta)
        Giter += gradients*gradients
        update = gradients*eta/(delta+np.sqrt(Giter))
        theta -= update
print("theta from own AdaGrad")
print(theta)

```

Running this code we note an almost perfect agreement with the results from matrix inversion.

RMSprop for adaptive learning rate with Stochastic Gradient Descent

```

# Using Autograd to calculate gradients using RMSprop and Stochastic Gradient descent
# OLS example
from random import random, seed
import numpy as np
import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

# Note change from previous example
def CostOLS(y,X,theta):
    return np.sum((y-X @ theta)**2)

n = 1000

```



```

x = np.random.rand(n,1)
y = 2.0+3*x +4*x*x# +np.random.randn(n,1)

X = np.c_[np.ones((n,1)), x, x*x]
XT_X = X.T @ X
theta_linreg = np.linalg.pinv(XT_X) @ (X.T @ y)
print("Own inversion")
print(theta_linreg)

# Note that we request the derivative wrt third argument (theta, 2 here)
training_gradient = grad(CostOLS,2)
# Define parameters for Stochastic Gradient Descent
n_epochs = 50
M = 5 #size of each minibatch
m = int(n/M) #number of minibatches
# Guess for unknown parameters theta
theta = np.random.randn(3,1)

# Value for learning rate
eta = 0.01
# Value for parameter rho
rho = 0.99
# Including AdaGrad parameter to avoid possible division by zero
delta = 1e-8
for epoch in range(n_epochs):
    Giter = 0.0
    for i in range(m):
        random_index = M*np.random.randint(m)
        xi = X[random_index:random_index+M]
        yi = y[random_index:random_index+M]
        gradients = (1.0/M)*training_gradient(yi, xi, theta)
        # Accumulated gradient
        # Scaling with rho the new and the previous results
        Giter = (rho*Giter+(1-rho)*gradients*gradients)
        # Taking the diagonal only and inverting
        update = gradients*eta/(delta+np.sqrt(Giter))
        # Hadamard product
        theta -= update
print("theta from own RMSprop")
print(theta)

```

And finally ADAM

```

# Using Autograd to calculate gradients using RMSprop and Stochastic Gradient descent
# OLS example
from random import random, seed
import numpy as np
import autograd.numpy as np
import matplotlib.pyplot as plt
from autograd import grad

# Note change from previous example
def CostOLS(y,X,theta):
    return np.sum((y-X @ theta)**2)

n = 1000
x = np.random.rand(n,1)
y = 2.0+3*x +4*x*x# +np.random.randn(n,1)

```

```

X = np.c_[np.ones((n,1)), x, x*x]
XT_X = X.T @ X
theta_linreg = np.linalg.pinv(XT_X) @ (X.T @ y)
print("Own inversion")
print(theta_linreg)

# Note that we request the derivative wrt third argument (theta, 2 here)
training_gradient = grad(CostOLS,2)
# Define parameters for Stochastic Gradient Descent
n_epochs = 50
M = 5 #size of each minibatch
m = int(n/M) #number of minibatches
# Guess for unknown parameters theta
theta = np.random.randn(3,1)

# Value for learning rate
eta = 0.01
# Value for parameters beta1 and beta2, see https://arxiv.org/abs/1412.6980
beta1 = 0.9
beta2 = 0.999
# Including AdaGrad parameter to avoid possible division by zero
delta = 1e-7
iter = 0
for epoch in range(n_epochs):
    first_moment = 0.0
    second_moment = 0.0
    iter += 1
    for i in range(m):
        random_index = M*np.random.randint(m)
        xi = X[random_index:random_index+M]
        yi = y[random_index:random_index+M]
        gradients = (1.0/M)*training_gradient(yi, xi, theta)
        # Computing moments first
        first_moment = beta1*first_moment + (1-beta1)*gradients
        second_moment = beta2*second_moment+(1-beta2)*gradients*gradients
        first_term = first_moment/(1.0-beta1**iter)
        second_term = second_moment/(1.0-beta2**iter)
        # Scaling with rho the new and the previous results
        update = eta*first_term/(np.sqrt(second_term)+delta)
        theta -= update
print("theta from own ADAM")
print(theta)

```

Introducing JAX

Presently, instead of using **autograd**, we recommend using **JAX**

JAX is Autograd and **XLA** (Accelerated Linear Algebra)), brought together for high-performance numerical computing and machine learning research. It provides composable transformations of Python+NumPy programs: differentiate, vectorize, parallelize, Just-In-Time compile to GPU/TPU, and more.

Here's a simple example on how you can use **JAX** to compute the derivate of the logistic function.

```

import jax.numpy as jnp
from jax import grad, jit, vmap

```

```

def sum_logistic(x):
    return jnp.sum(1.0 / (1.0 + jnp.exp(-x)))

x_small = jnp.arange(3.)
derivative_fn = grad(sum_logistic)
print(derivative_fn(x_small))

```

Computing quantum gradients

After this overview over various gradient approaches, let us implement efficient implementations of gradient methods to the derivatives of the Hamiltonian expectation values.

```

from matplotlib import pyplot as plt
import numpy as np
from scipy.optimize import minimize
dim = 2
Hamiltonian = np.zeros((dim,dim))
e0 = 0.0
e1 = 4.0
Xnondiag = 0.20
Xdiag = 3.0
Eigenvalue = np.zeros(dim)
# setting up the Hamiltonian
Hamiltonian[0,0] = Xdiag+e0
Hamiltonian[0,1] = Xnondiag
Hamiltonian[1,0] = Hamiltonian[0,1]
Hamiltonian[1,1] = e1-Xdiag
# diagonalize and obtain eigenvalues, not necessarily sorted
EigValues, EigVectors = np.linalg.eig(Hamiltonian)
permute = EigValues.argsort()
EigValues = EigValues[permute]
# print only the lowest eigenvalue
print(EigValues[0])

# Now rewrite it in terms of the identity matrix and the Pauli matrix X and Z
X = np.array([[0,1],[1,0]])
Y = np.array([[0,-1j],[1j,0]])
Z = np.array([[1,0],[0,-1]])
# identity matrix
I = np.array([[1,0],[0,1]])

epsilon = (e0+e1)*0.5; omega = (e0-e1)*0.5
c = 0.0; omega_z=Xdiag; omega_x = Xnondiag
Hamiltonian = (epsilon+c)*I+(omega_z+omega)*Z+omega_x*X
EigValues, EigVectors = np.linalg.eig(Hamiltonian)
permute = EigValues.argsort()
EigValues = EigValues[permute]
# print only the lowest eigenvalue
print(EigValues[0])

# define the rotation matrices

def Rx(theta):
    return np.cos(theta*0.5)*I-1j*np.sin(theta*0.5)*X
def Ry(phi):

```

```

        return np.cos(phi*0.5)*I-1j*np.sin(phi*0.5)*Y

#define basis states
basis0 = np.array([1,0])
basis1 = np.array([0,1])

# Computing the expectation value of the energy
def Energy(theta,phi):
    Basis = Ry(phi) @ Rx(theta) @ basis0
    energy = Basis.conj().T @ Hamiltonian @ Basis
    return energy

# Set up iteration using gradient descent method
eta = 0.1
Niterations = 100
# Random angles using uniform distribution
theta = 2*np.pi*np.random.rand()
phi = 2*np.pi*np.random.rand()
pi2 = 0.5*np.pi
for iter in range(Niterations):
    thetagradient = 0.5*(Energy(theta+pi2,phi)-Energy(theta-pi2,phi))
    phigradient = 0.5*(Energy(theta,phi+pi2)-Energy(theta,phi-pi2))
    theta -= eta*thetagradient
    phi -= eta*phigradient
print(Energy(theta,phi))

```

A smarter way of doing this

The above approach means that we are setting up several matrix-matrix and matrix-vector multiplications. Although straight forward it is not the most efficient way of doing this, in particular in case the matrices become large (and sparse). But there are some more important issues.

In a physical realization of these systems we cannot just multiply the state with the Hamiltonian. When performing a measurement we can only measure in one particular direction. For the computational basis states which we have, $|0\rangle$ and $|1\rangle$, we have to measure along the bases of the Pauli matrices and reconstruct the eigenvalues from these measurements.

From our earlier discussions we know that the Pauli Z matrix has the above basis states as eigen states through

$$\sigma_z|0\rangle = Z|0\rangle = +1|0\rangle,$$

and

$$\sigma_z|1\rangle = Z|1\rangle = -1|1\rangle,$$

with eigenvalue -1 .

For the Pauli X matrix on the other hand we have

$$\sigma_x|0\rangle = X|0\rangle = +1|1\rangle,$$

and

$$\sigma_x|1\rangle = X|1\rangle = -1|0\rangle,$$

with eigenvalues 1 in both cases. The latter two equations tell us that the computational basis we have chosen, and in which we will prepare our states, is not an eigenbasis of the σ_x matrix.

We will thus try to rewrite the Pauli X matrix in terms of a Pauli Z matrix. Fortunately this can be done using the Hadamard matrix twice, that is

$$\mathbf{X} = \sigma_x = \mathbf{H}\mathbf{Z}\mathbf{H}.$$

The Pauli Y matrix can be written as

$$\mathbf{Y} = \sigma_y = \mathbf{H}\mathbf{S}^\dagger\mathbf{Z}\mathbf{H}\mathbf{S},$$

where S is the phase matrix

$$S = \begin{bmatrix} 1 & 0 \\ 0 & i \end{bmatrix}.$$

From here and on we will denote the Pauli matrices by X , Y and Z and we can write the expectation value of the Hamiltonian as

$$\langle\psi|(c + \mathcal{E})\mathbf{I} + (\Omega + \omega_z)\mathbf{Z} + \omega_x\mathbf{H}\mathbf{Z}\mathbf{H}|\psi\rangle,$$

which we can rewrite as

$$(c + \mathcal{E})\langle\psi|\mathbf{I}|\psi\rangle + (\Omega + \omega_z)\langle\psi|\mathbf{Z}|\psi\rangle + \omega_x\langle\psi\mathbf{H}|\mathbf{Z}|\mathbf{H}\psi\rangle.$$

The first and second term are too easy to perform a measurement on since we just need to compute $\langle\psi|\mathbf{I}|\psi\rangle$ and $\langle\psi|\mathbf{Z}|\psi\rangle$. For the final term we need just to add the action of the Hadamard matrix and we are done.

To do: Set up codes for this using gradient descent and perform a series of measurements

Two-qubit system and the VQE

We extend now the system to a two-qubit system with the following computational basis states and Hamiltonian matrix written out in terms of Pauli spin matrices.

This system can be thought of as composed of two subsystems A and B . Each subsystem has computational basis states

$$|0\rangle_{A,B} = \begin{bmatrix} 1 & 0 \end{bmatrix}^T \quad |1\rangle_{A,B} = \begin{bmatrix} 0 & 1 \end{bmatrix}^T.$$

The subsystems could represent single particles or composite many-particle systems of a given symmetry. This leads to the many-body computational basis states

$$|00\rangle = |0\rangle_A \otimes |0\rangle_B = \begin{bmatrix} 1 & 0 & 0 & 0 \end{bmatrix}^T,$$

and

$$|01\rangle = |0\rangle_A \otimes |1\rangle_B = \begin{bmatrix} 0 & 1 & 0 & 0 \end{bmatrix}^T,$$

and

$$|10\rangle = |1\rangle_A \otimes |0\rangle_B = \begin{bmatrix} 0 & 0 & 1 & 0 \end{bmatrix}^T,$$

and finally

$$|11\rangle = |1\rangle_A \otimes |1\rangle_B = \begin{bmatrix} 0 & 0 & 0 & 1 \end{bmatrix}^T.$$

These computational basis states define also the eigenstates of the non-interacting Hamiltonian

$$H_0|00\rangle = \epsilon_{00}|00\rangle,$$

$$H_0|10\rangle = \epsilon_{10}|10\rangle,$$

$$H_0|01\rangle = \epsilon_{01}|01\rangle,$$

and

$$H_0|11\rangle = \epsilon_{11}|11\rangle.$$

The interacting part of the Hamiltonian H_I is given by the tensor product of two σ_x and σ_z matrices, respectively, that is

$$H_I = H_x \sigma_x \otimes \sigma_x + H_z \sigma_z \otimes \sigma_z,$$

where H_x and H_z are interaction strength parameters. Our final Hamiltonian matrix is given by

$$\mathbf{H} = \begin{bmatrix} \epsilon_{00} + H_z & 0 & 0 & H_x \\ 0 & \epsilon_{10} - H_z & H_x & 0 \\ 0 & H_x & \epsilon_{01} - H_z & 0 \\ H_x & 0 & 0 & \epsilon_{11} + H_z \end{bmatrix}.$$

Switching to Qiskit (notes developed by Stian Bilek)

```
import numpy as np
import qiskit as qk
from scipy.optimize import minimize

# Initialize registers and circuit
n_qubits = 1 #Number of qubits
n_cbits = 1 #Number of classical bits (the number of qubits you want to measure at the end of the
qreg = qk.QuantumRegister(n_qubits) #Create a quantum register
creg = qk.ClassicalRegister(n_cbits) #Create a classical register
circuit = qk.QuantumCircuit(qreg,creg) #Create your quantum circuit
circuit.draw() #Draw circuit. It is empty
# Perform operations on qubit
circuit.x(qreg[0]) #Applies a Pauli X gate to the first qubit in the quantum register
circuit.draw()
# Chose a qubit to measure and encode the results to a classical bit
#Measure the first qubit in the quantum register
#and encode the results to the first qubit in the classical register
circuit.measure(qreg[0],creg[0])
circuit.draw()
# Execute circuit

backend = qk.Aer.get_backend('qasm_simulator')
#This is the device you want to use. It is an ideal simulation of a quantum device
```

```

job = backend.run(circuit,shots=1000) #Run the circuit 1000 times
result = job.result()
counts = result.get_counts()
print(counts)
circuit.clear()
circuit.draw()

circuit.h(qreg[0]) #Apply a Hadamard gate to the first qubit of the quantum register
circuit.measure(qreg,creg)
print(circuit.draw())
job = backend.run(circuit,shots=1000)
result = job.result()
counts = result.get_counts()
print(counts)
circuit.clear()

```

Create a two-qubit circuit and set up a Bell state

```

n_qubits = 2
n_cbits = 2
qreg = qk.QuantumRegister(n_qubits)
creg = qk.ClassicalRegister(n_cbits)
circuit = qk.QuantumCircuit(qreg,creg)
circuit.draw()
circuit.h(qreg[0])
circuit.cx(qreg[0],qreg[1])

```

This is a controlled operation. Apply a Pauli X gate to the second qubit ($qreg[1]$) if the first qubit ($qreg[0]$) is in the $|1\rangle$ state. Else do nothing

```

circuit.draw()
circuit.measure(qreg,creg)
circuit.draw()
job = backend.run(circuit,shots=1000)
result = job.result()
counts = result.get_counts()
print(counts)
circuit.clear()

```

Apply rotation to qubit

```

theta = np.pi/3
circuit.rx(theta, qreg[0]) #R_x(theta) rotation on the first qubit (qreg[0])
circuit.measure(qreg,creg)
print(circuit.draw())
job = backend.run(circuit,shots=1000)
result = job.result()
counts = result.get_counts()
circuit.clear()
print(counts)

```

Find the lowest eigenvalue of

$$H = c_1 Z_0 + c_2 Z_1 + c_3 X_0 Y_1$$

We will use

$$\langle \psi | H | \psi \rangle = c_1 \langle \psi | Z_0 | \psi \rangle + c_2 \langle \psi | Z_1 | \psi \rangle + c_3 \langle \psi | X_0 Y_1 | \psi \rangle$$

```
I = np.eye(2)
X = np.array([[0,1],[1,0]])
Y = np.array([[0,-1j],[1j,0]])
Z = np.array([[1,0],[0,-1]])
H = np.kron(Z,I) + np.kron(I,Z) + np.kron(X,Y)
eigvals,eigvecs = np.linalg.eigh(H)
print(eigvals[0])

c_1 = 1
c_2 = 1
c_3 = 1

h_1 = [c_1,[0],['z']]
h_2 = [c_2,[1],['z']]
h_3 = [c_3,[0,1],['x','x']]
H = [h_1,h_2,h_3]
```

Create ansatz

```
def ansatz(theta,n_qubits):
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)
    for i in range(n_qubits):
        circuit.ry(theta[i],qreg[i])
    for i in range(n_qubits-1):
        circuit.cx(qreg[i],qreg[i+1])
    return(circuit)
qreg = qk.QuantumRegister(n_qubits)
circuit = qk.QuantumCircuit(qreg)
circuit.h(qreg[:2])
print('Before ansatz')
print(circuit.draw())
theta = np.random.randn(2)
n_qubits = 2
circuit = circuit.compose(ansatz(theta,n_qubits))
print('After ansatz')
circuit.draw()
```

Change measurement basis

```
def basis_change(h_i,n_qubits):
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)

    for qubit,operator in zip(h_i[1],h_i[2]):
        if operator == 'x':
            circuit.h(qreg[qubit])
        if operator == 'y':
            circuit.sdg(qreg[qubit])
            circuit.h(qreg[qubit])
    return(circuit)
n_qubits = 2
qreg = qk.QuantumRegister(n_qubits)
```



```

circuit = qk.QuantumCircuit(qreg)
theta = np.random.randn(n_qubits)
circuit = circuit.compose(ansatz(theta,n_qubits))
print('Ansatz circuit')
circuit.draw()
circuit = circuit.compose(basis_change(H[2],n_qubits))
print('After basis transformation:')
print(circuit.draw())

```

Get energy for given rotational parameters, theta

```

def get_energy(theta):
    n_qubits = 2
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)
    circuit = circuit.compose(ansatz(theta,n_qubits))
    circuit_list = []
    for idx,h_i in enumerate(H):
        basis_change_circuit = basis_change(h_i,n_qubits)
        new_circuit = circuit.compose(basis_change_circuit)
        creg = qk.ClassicalRegister(len(h_i[1]))
        new_circuit.add_register(creg)
        new_circuit.measure(qreg[h_i[1]],creg)
        circuit_list.append(new_circuit)
    shots = 10000
    job = backend.run(circuit_list,shots=shots)
    E = np.zeros(len(circuit_list))
    for i in range(len(circuit_list)):
        result = job.result()
        counts = result.get_counts(i)
        for key,value in counts.items():
            e = 1
            for bit in key:
                if bit == '0':
                    e *= 1
                if bit == '1':
                    e *= -1
            E[i] += e*value
        E[i] *= H[i][0]
    E /= shots
    return(np.sum(E))

theta = np.random.randn(2)
get_energy(theta)

```

Minimize energy with Scipy

```

theta = np.random.randn(2)
res = minimize(get_energy, theta, method='Powell',tol=1e-12)
get_energy(res.x)

```

We might need a more flexible ansatz

```

def ansatz(theta,n_qubits):
    qreg = qk.QuantumRegister(n_qubits)
    circuit = qk.QuantumCircuit(qreg)

```

```

idx = 0
for i in range(n_qubits):
    circuit.ry(theta[idx],qreg[i])
    idx += 1
for i in range(n_qubits-1):
    circuit.cx(qreg[i],qreg[i+1])
for i in range(n_qubits):
    circuit.rx(theta[idx],qreg[i])
    idx += 1
for i in range(n_qubits-1):
    circuit.cx(qreg[i],qreg[i+1])
return(circuit)
theta = np.random.randn(4)
res = minimize(get_energy, theta, method='Powell',tol=1e-16)
get_energy(res.x)

```

Then we minimize energy with gradient descent

$$\frac{\partial E(\theta_1, \dots, \theta_i, \dots, \theta_p)}{\partial \theta_i} = \frac{E(\theta_1, \dots, \theta_i + \pi/2, \dots, \theta_p) - E(\theta_1, \dots, \theta_i - \pi/2, \dots, \theta_p)}{2}.$$

```

epochs = 200
theta = np.random.randn(4)
for epoch in range(epochs):
    print(epoch,get_energy(theta))
    grad = np.zeros_like(theta)
    for idx in range(theta.shape[0]):
        theta_temp = theta.copy()
        theta_temp[idx] += np.pi/2
        E_plus = get_energy(theta_temp)
        theta_temp[idx] -= np.pi
        E_minus = get_energy(theta_temp)
        grad[idx] = (E_plus - E_minus)/2
    theta -= 0.1*grad

```

Moving to the Lipkin model

To construct an efficient ansatz, we must determine the subspace within which the Hamiltonian lives. To begin, note that particles are only ever moved between energy levels in pairs. This implies that all possible states have a Hamming weight of constant parity (odd or even); this is the same as the signature r being conserved. Further, note that the Hamiltonian's coefficients (ϵ and V) are state independent (do not depend on the indices n or m) as the states labeled by these indices are degenerate and thus have the same energy level. Thus, the Hamiltonian treats all states with the same number of excited particles (Hamming weight of the state) as the same. Therefore, the following ansatz forms exactly cover the subspace within which the N -degenerate Hamiltonian explores:

$$|\psi_{\text{even}}\rangle = \sum_{k=0}^{\lfloor n/2 \rfloor} c_{2k} |D_{2k}^n\rangle, \quad (68)$$

$$|\psi_{\text{odd}}\rangle = \sum_{k=0}^{\lfloor n/2 \rfloor} c_{2k+1} |D_{2k+1}^n\rangle. \quad (69)$$

Here $|D_k^n\rangle$ represents a Dicke state which is defined as equal superposition of all n -qubit states with Hamming weight k . That is

$$|D_k^n\rangle = \frac{1}{\sqrt{\binom{n}{k}}} \sum_{x \in h_k^n} |x\rangle, \quad (70)$$

where $h_k^n = \{|x\rangle \mid l(x) = n, \text{wt}(x) = k\}$. There are two ways we can think of to prepare such ansatz: The first is to prepare them exactly as it is known how to deterministically prepare Dicke states with linear depth. The reference provides an algorithm for preparing a set of gates U_k^n that prepares a Dicke state from a product state of Hamming weight k ; that is

$$U_k^n |1\rangle^{\otimes k} |0\rangle^{\otimes n-k} = |D_k^n\rangle. \quad (71)$$

It then describes how to one can create an arbitrary superposition of Dicke states, which we modify here to restrict ourselves to a Hamming weight of constant parity. The circuit to construct such a state (for the $k = 6$ case, as an example) is given below

```
import numpy as np
import qiskit
from qiskit.visualization import circuit_drawer
from qiskit.quantum_info import Statevector
from matplotlib.pyplot import figure
from qiskit import QuantumRegister, QuantumCircuit, ClassicalRegister, Aer, assemble
from qiskit.providers.aer.noise import NoiseModel
import pylatexenc
from qiskit.algorithms import VQE
from qiskit.utils import QuantumInstance
from qiskit.opflow import X, Z, I, Y
from qiskit.circuit import Parameter
from qiskit.algorithms.optimizers import ADAM
from qiskit.opflow import AerPauliExpectation
from qiskit import IBMQ
import cmath
import pandas as pd
from scipy.sparse import diags
import numpy.linalg as LA
import matplotlib.pyplot as plt
from IPython.display import Image
import warnings
warnings.filterwarnings('ignore')
pi=np.pi
```

```

#function that sorts eigenvalues with its eigenvectors in ascending order
def eigen(A):
    eigenValues, eigenVectors = LA.eig(A)
    idx = np.argsort(eigenValues)
    eigenValues = eigenValues[idx]
    eigenVectors = eigenVectors[:,idx]
    return (eigenValues, eigenVectors)

#one body expectation value
def one_body(E,N):
    k = N/2
    m = np.arange(-k,k+1,1) # Since the collective space is Omega+1
    return E*np.diag(m) #return a matrix where its diagonal elements are epsilon*K_0

#two body expectation value
def two_body(V,N):
    k = N/2
    m = np.arange(-k,k+1,1)
    left = np.zeros(len(m)-2,dtype=complex)
    right = np.zeros(len(m)-2,dtype=complex)
    diag = np.zeros(len(m),dtype=complex)
    for i in range(len(left)):
        CG = cmath.sqrt(k*(k+1)-(m[i]+2)*(m[i]+1))*cmath.sqrt(k*(k+1)-m[i]*(m[i]+1)) #calculate Clebs
        left[i] = CG
        right[i] = CG
    k = [left,diag,right]
    offset = [-2,0,2]
    return -0.5*V*diags(k,offset).toarray() #return a matrix where its off diagonal elements are (1/

#full expectation value
def quasi_spin(E,V,N):
    ob = one_body(E,N)
    tb = two_body(V,N)
    H = ob+tb
    e,v = eigen(H) # find the eigenvalues of the Hamiltonian
    return e,H

#converts chi to V
def Vp(E,omega,chi):
    return (chi*E)/(omega-1)

#parameters
E = 1
chi = np.arange(0,2.1,0.1)
omega = 2

EVO = []
EV1 = []
EV2 = []
Ham = []
for i in chi:
    v = Vp(E,omega,i)
    EigenV,H = quasi_spin(E,v,omega) #return eigenvalues and Hamiltonian
    Ham.append(H)
    EVO.append(EigenV[0])
    EV1.append(EigenV[1])
    EV2.append(EigenV[2])

```

```

matrix = pd.DataFrame(Ham[5].real)
print('Hamiltonian matrix')
matrix.head()

plt.plot(chi,EV0)
plt.plot(chi,EV1)
plt.plot(chi,EV2)
plt.xlabel('$\chi$')
plt.ylabel('Energy')
plt.title('$\Omega=2$ exact Lipkin Model')

def LM_circuit():
    theta = Parameter('theta')
    QC = QuantumCircuit(2)
    QC.ry(2*(theta-np.pi/2),0)
    QC.cnot(0,1)
    return QC

QC = LM_circuit()
QC.draw(output='mpl')

```

VQE method. For this method, we will be using Qiskit's VQE function, where we specify

- Quantum circuit
 1. Optimizer
 2. Quantum instance (i.e. which backend). Here we will be using the "qasm_ssimulator"
- 3. Initial point (i.e. θ search space)
- 4. Hamiltonian/measurement basis

To define the Hamiltonian, we will use the Qiskit Pauli operator functions I,Z,X,Y. In this method, for a given χ , we will do a search over θ from $-\frac{\pi}{2}$ to $\frac{\pi}{2}$ and picking out the minimum energy value

```

#sim1 min example
sim = Aer.get_backend('qasm_simulator')
adam =qiskit.algorithms.optimizers.ADAM(maxiter=10000) #optimizer for VQE
epsilon = 1
omega = 2
chi2 = [0.5,1]

t = np.arange(-pi/2,pi/2,0.05) #0.1 step size finishes pretty fast
energy = []
for x in chi2:
    ev = []
    v = Vp(epsilon,omega,x)
    for i in range(len(t)):
        H = 0.5 * epsilon * ( Z ^ I ) + 0.5 * epsilon * ( I ^ Z ) -0.5 * v * ( X ^ X ) +0.5 * v *
        vqe = VQE(ansatz=LM_circuit(),optimizer=adam,initial_point=t[i],quantum_instance=sim,exp

```

```

        result = vqe.compute_minimum_eigenvalue(H)
        ev.append(result.eigenvalue)
    energy.append(ev)

plt.plot(t,energy[0],label='$\chi=0.5$')
plt.plot(t,energy[1],label='$\chi=1$')
plt.xlabel('$\theta$')
plt.ylabel('Ground State Energy')
plt.title('$E_{g.s}$ vs. $\theta$')
plt.legend()

epsilon = 1
omega = 2
chi3 = np.arange(0,2.1,0.1)
t = np.arange(-pi/2,pi/2,0.1)
adam =qiskit.algorithms.optimizers.ADAM(maxiter=10000)
LM_sim = []
for x in chi3:
    v = Vp(epsilon,omega,x)
    ev = []
    for i in range(len(t)):
        H = 0.5 * epsilon * ( Z ^ I ) +0.5 * epsilon * ( I ^ Z ) -0.5 * v * ( X ^ X ) +0.5 * v * (
        vqe = VQE(ansatz=LM_circuit(),optimizer=adam,initial_point=[t[i]],quantum_instance=sim,exp
        result = vqe.compute_minimum_eigenvalue(H)

        ev.append(result.eigenvalue)
    LM_sim.append(min(ev))

plt.plot(chi,np.array(exact),label='$E_{exact}$',color='b')
plt.plot(chi,np.array(HF),label='$E_{HF}$',color='g')
plt.scatter(chi3,LM_sim,label='QC sim',color='cyan')
plt.axvline(1, color = 'k', linestyle='--')
plt.ylabel('$E_{g.s}$')
plt.xlabel('$\chi$')
plt.title(f'$\Omega=2$')
plt.legend()

```

VQE method using Qiskit. For this method, we will be using Qiskit's VQE function, where we specify

1. Quantum circuit
2. Optimizer
3. Quantum instance (i.e. which backend). Here we will be using the "statevector_simulator"
4. Hamiltonian/measurement basis

```

epsilon = 1
omega = 2
chi5 = np.arange(0,2.1,0.1)
adam =qiskit.algorithms.optimizers.ADAM(maxiter=10000)
LM_sim = []
for x in chi5:
    v = Vp(epsilon,omega,x)

```

```

H = 0.5 * epsilon * ( Z ^ I ) + 0.5 * epsilon * ( I ^ Z ) - 0.5 * v * ( X ^ X ) + 0.5 * v * ( Y ^ Y )
vec = Aer.get_backend('statevector_simulator')
vqe = VQE(ansatz=LM_circuit(), optimizer=adam, quantum_instance=vec, expectation=AerPauliExpectation())
result = vqe.compute_minimum_eigenvalue(H)
LM_sim.append(result.eigenvalue)

plt.plot(chi, np.array(exact), label='$E_{exact}$', color='b')
plt.plot(chi, np.array(HF), label='$E_{HF}$', color='g')
plt.scatter(chi5, LM_sim, label='QC state vec', color='cyan')
plt.axvline(1, color='k', linestyle='--')
plt.ylabel('$E_{g.s}$')
plt.xlabel('$\chi$')
plt.title('$\Omega=2$')
plt.legend()

```